

ECE428 MT1 cheat-sheet

Latency: delay (time taken) for a bit transmission

Bandwidth: Total amount of info. transmitted over the channel per unit time

Clock skew / Clock drift (rate)

Asynchronous system:

No bound on process execution speeds, message passing delay, clock drift rate.

Periodic ping and ack time out $\Delta_1 = 2 \text{ (max delay)}$

Periodic heartbeats $T + \Delta_2$, $\Delta_2 = (\text{Max} - \text{min}) \text{ delay}$

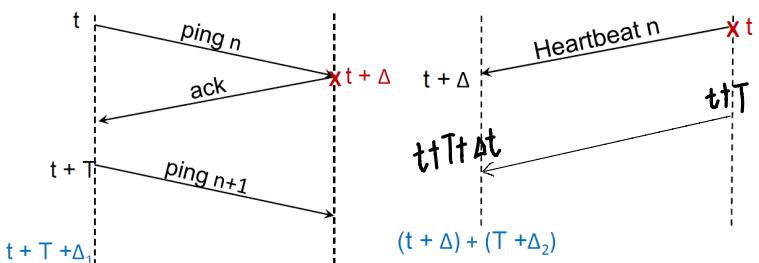
Failure Detection: Completeness (eventually) $\Theta(\text{Centralized})$

Accuracy (failure \rightarrow the crashed process) $\Theta(\text{Ring hb})$ $\Theta(\text{all-to-all})$

Worst case failure detection time: T_{fd}

Ping-ack: $T + \Delta_1 - \Delta$ Heartbeats: $T + \Delta_2 + \Delta$

T : Time period Δ : network delay / Time between two processes



$T \downarrow$, $T_{fd} \downarrow$, $B \uparrow$; $\Delta_1, \Delta_2 \uparrow$, accuracy \uparrow , $T_{fd} \uparrow$

Communication
Omission: process/channel crash \rightarrow failed msg transmission

Arbitrary (Byzantine) failure: executing wrongly \rightarrow wrong msg

Time failure: Timing guarantee are not met \rightarrow fail stop

Sync systems have bound on max drift rate.

Clocks in a system externally synchronized within bound $D \Rightarrow$ internally sync within a bound of $2D$.

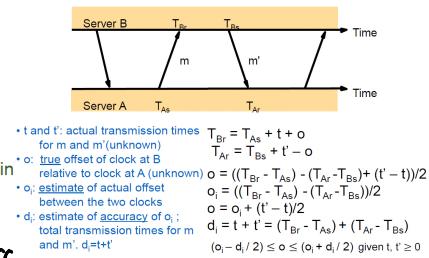
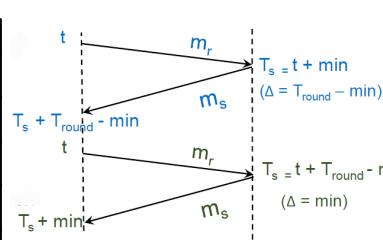
Sync:

$T_c = T_s + \frac{\text{min} + \text{max}}{2} \text{ delay} / 2 \Rightarrow$ skew(client, sever) $\leq \frac{(\text{max} - \text{min})}{2} \text{ delay}$

Cristian Algorithm: RTT $\leq \frac{(\text{max} - \text{min})}{2} \text{ delay}$

ASync: $T_c = T_s + (T_{\text{round}}/2) \Rightarrow$ skew $\leq (T_{\text{round}}/2) - \text{min delay}$

Estimates Δ as RTT/2



Cannot handle faulty t servers

Berkeley Algorithm: Average all local t and send back offset (Support internal sync.)

Network Time Protocol (NTP)

Hierarchical structure for scalability

Multiple lower strata servers for robustness (LAN)

Authentication mechanisms for security

Statistical techniques for better accuracy (Procedure-call & symmetric mode)

Handed-before: (HB) $e \rightarrow e'$

$e \rightarrow_i e'$: e happened before e' , as observed by p_i

$a \rightarrow e$ and $e \rightarrow a \Rightarrow a \parallel e$, a and e are concurrent

Lamport's Logical clock HB ✓

$e \rightarrow e' \Rightarrow L(e) < L(e')$ $L(e) < L(e') \Rightarrow e \rightarrow e'$ or $e \parallel e'$

Vector Timestamps

$e \rightarrow e'$ iff $V(e) < V(e')$ $e \parallel e'$ iff $(V \neq V' \text{ and } V' \neq V)$

Global State (Global Snapshot)

For a process p_i , history $(p_i) = h_i = \langle e_i^0, e_i^1, \dots \rangle$ prefix his

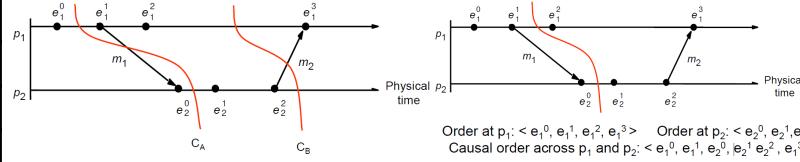
s_i^k : p_i 's state immediately after k th event. $h_i^k = \langle e_i^0, e_i^1, \dots, e_i^k \rangle$

Global history $H = \bigcup_i (h_i)$ Global state $S = \bigcup_i (s_i)$

a cut $C \subseteq H = h_1^{c_1} \cup h_2^{c_2} \cup \dots \cup h_n^{c_n}$

the frontier of $C = \{e_i^{c_i}, i=1, 2, \dots, n\}$

global state S corresponds to cut $C = \bigcup_i (s_i^{c_i})$



$C_A: \langle e_1^0, e_1^2 \rangle$
Frontier of $C_A: \langle e_1^0, e_1^2 \rangle$
Cut Inconsistent cut.

$C_B: \langle e_1^0, e_1^1, e_1^2, e_2^0, e_2^1, e_2^2 \rangle$
Frontier of $C_B: \langle e_1^2, e_2^2 \rangle$
Consistent cut.

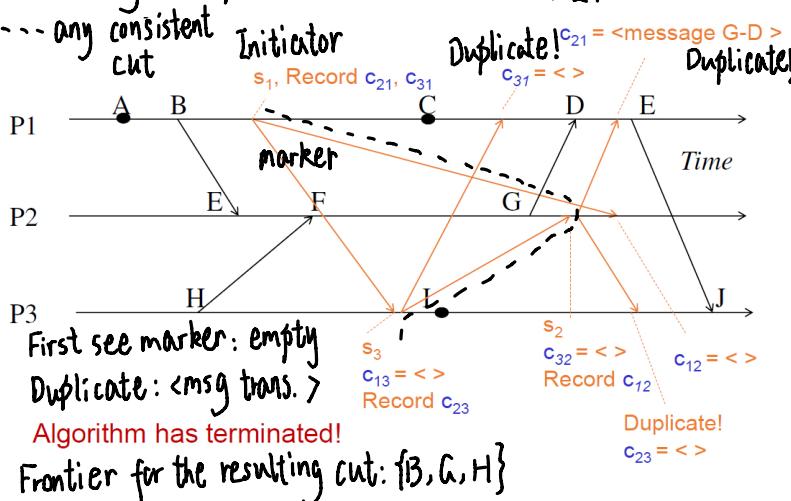
Consistent iff $\forall e \in C, \text{ if } f \rightarrow e,$
then $f \in C$

Global state S consistent
iff corresponding to cons. cut.

Order at $p_1: < e_1^0, e_1^1, e_1^2, e_2^0, e_2^1, e_2^2 >$
Order at $p_2: < e_2^0, e_2^1, e_2^2 >$
Causal order across p_1 and $p_2: < e_1^0, e_1^1, e_2^0, e_1^2, e_2^1, e_2^2, e_1^3 >$

Linearization: $< e_1^0, e_1^1, e_1^2, e_2^0, e_2^1, e_2^2, e_1^3 >$
Linearization: $< e_1^0, e_1^1, e_2^0, e_1^2, e_2^1, e_2^2, e_1^3 >$

Chandy-Lamport Algorithm: pair communication



Execution Lattice: Each path represents a linearization

Liveness: iff $\text{liveness}(P(S_0)) \equiv \forall L \in \text{linearizations from } S_0$

S_0 . L passes through a SL & $P(SL) = \text{True}$

② For all linearizations starting from S_0 , P is True for some state S_L reachable from S_0 . i.e. terminate; fail detect complete

Safety: iff ① safety($P(S_0)$) $\equiv \forall s$ reachable from S_0 , $P(s) = \text{true}$ ② For all states s reachable from S_0 , $P(s)$ is True

Unicast: one to one; Broadcast: one to all; Multicast: any one to a group
Basic (B-Multicast): each in unicast.

Reliable (R-Multicast):

Validity and Agreement ensure liveness: All correct P (itself m) others agree on m) deliver msg m.

FIFO ordering: focus on the delivering order of msg sent from one process

Causal ordering: Lamport's HB rule

Total ordering (Atomic Broadcast):

Ensure all the processes deliver multicast in the same order

Implement FIFO ordering: sequencer (pre-sender) FIFO

Implement causal ordering: buffering until causality

- FIFO ordering

- If a correct process issues multicast(g, m) and then multicast(g, m'), then every correct process that delivers m' will have already delivered m . M1:1, M1:2...

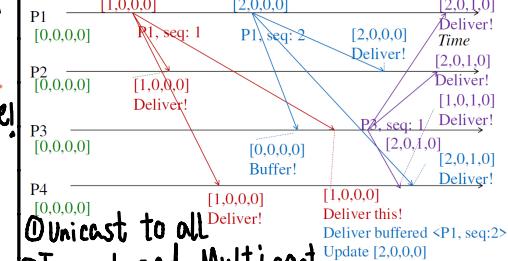
- Causal ordering

- If multicast(g, m) \rightarrow multicast(g, m') then any correct process that delivers m' will have already delivered m .
- Note that \rightarrow counts multicast messages delivered to the application, rather than all network messages. M1:1, M2:1...

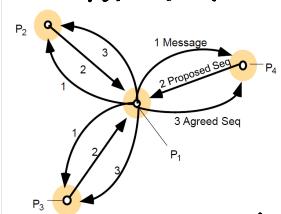
- Total ordering

- If a correct process delivers message m before m' (independent of sending order), then any other correct process that delivers m' will have already delivered m . M1:1, M2:1, M3:1, M3:2

FIFO Execution



@Decen. ISIS Queue



① Central server based algorithm

Breaking deadlocks (MA in mutual)

- System of 6 processes {0, 1, 2, 3, 4, 5}: 0, 1, 2 want to enter critical section:
 - V₀ = {0, 1, 2}: 0, 2 send reply to 0, but 1 sends reply to 1;
 - V₁ = {1, 3, 5}: 1, 3 send reply to 1, but 5 sends reply to 2;
 - V₂ = {2, 4, 5}: 4, 5 send reply to 2, but 2 sends reply to 0;
- Suppose (L1, P1) < (L0, P0) < (L2, P2).
- P2 will send fail to itself when it receives its own request after P0.
- P5 will send inquire to P2 when it receives P1's request.
- P2 will send relinquish to V₂. P5 and P4 will set "voted = false". P5 will reply to P1.
- P1 can now enter CS, followed by P0, and then P2.

Mutual Exclusion

enter(), AccessResource(), exit()

i.e. wait(s) critical section Problem i.e. signal(s)

1. Central Server Algorithm (hotspot on the leader)

Leader keeps a queue of waiting requests from processes that require to access CS, a token that allows holder to access CS. safety and liveness ✓, request + guarantee client delay is focused on each enter. (exit too but ignored)
Synchronization delay is the T between the pre-exit and the new enter. (throughput of the system) measure of the release + guarantee Bandwidth: 2 msg enter + 1 msg exist

2. Ring-based Algorithm

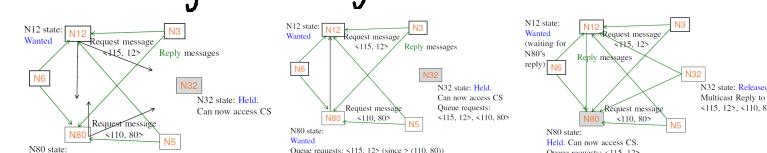
Bandwidth: 1 msg / N msg per enter for the req. P / the system
1 msg exist

Client delay: 0 to N (just receive/send out token) O(N)

Synchronization delay: 1 to (N-1) O(N)

process to enter is: successor predecessor of the Pexit

3. Ricart-Agrawala's Algorithm



Held until receive all replies; reply by FIFO priority no token. Worst-case: smaller T in timestamp

wait IN-1 P to send Reply

Bandwidth: 2*(N-1) per enter, (N-1) per exist

Client delay: one round-trip time (RTT) O(1)

Synchronization delay: one msg transmission time. O(1)

4. Maekawa's Algorithm: voting sets by only some G members

\sqrt{N} by FN matrix works best. Bandwidth: $2\sqrt{N} - 1$ O(\sqrt{N})

client delay: 1 RTT; synchronization delay: 2 msg trans. T safety; Liveness: A deadlock could occur; Orderx