

Distributed Systems

ECE428

Lecture 8

Adopted from Spring 2021

Today's agenda

- Multicast
 - Chapter 15.4
- Goal: reason about desirable properties for message delivery among a group of processes.

Recap: Multicast

- Useful communication mode in distributed systems:
 - Writing an object across replica servers.
 - Group messaging.
 -
- Basic multicast (B-multicast): unicast send to each process in the group.
 - Does not guarantee consistent message delivery if sender fails.
- Reliable multicast (R-multicast):
 - Defined by three properties: *integrity*, *validity*, *agreement*.
 - If some correct process multicasts a message *m*, then all other correct processes deliver the *m* (exactly once).
 - When a process receives a message 'm' for the first time, it re-multicasts it again to other processes in the group.

Recap: Ordered Multicast

- FIFO ordering
 - If a correct process issues $\text{multicast}(g, m)$ and then $\text{multicast}(g, m')$, then every correct process that delivers m' will have already delivered m .
- Causal ordering
 - If $\text{multicast}(g, m) \rightarrow \text{multicast}(g, m')$ then any correct process that delivers m' will have already delivered m .
 - Note that \rightarrow counts multicast messages delivered to the application, rather than all network messages.
- Total ordering
 - Yet to discuss.

Recap: Ordered Multicast

- FIFO ordering
 - If a correct process issues $\text{multicast}(g, m)$ and then $\text{multicast}(g, m')$, then every correct process that delivers m' will have already delivered m .
- Causal ordering
 - If $\text{multicast}(g, m) \rightarrow \text{multicast}(g, m')$ then any correct process that delivers m' will have already delivered m .
 - Note that \rightarrow counts multicast messages delivered to the application, rather than all network messages.
- Total ordering
 - Yet to discuss.


Where is causal ordering useful?

- Group = set of your friends on a social network.
- A friend sees your message m , and she posts a response (comment) m' to it.
 - If friends receive m' before m , it wouldn't make sense
 - But if two friends post messages m'' and n'' concurrently, then they can be seen in any order at receivers.
- A variety of systems implement causal ordering:
 - social networks, bulletin boards, comments on websites, etc.

Causal vs FIFO


- Causal Ordering \Rightarrow FIFO Ordering
- Why?
 - If two multicasts M and M' are sent by the same process P , and M was sent before M' , then $M \rightarrow M'$.
 - Then a multicast protocol that implements causal ordering will obey FIFO ordering since $M \rightarrow M'$.
- Reverse is not true! FIFO ordering does not imply causal ordering.

Example



Does this satisfy FIFO order?


Example



Does this satisfy FIFO order?


Yes

Example



Does this satisfy causal order?


Example



Does this satisfy causal order?

No


Example



M1:1 is delivered at P3 after M3:1's multicast.
Does this satisfy causal order?

Yes


Example



Does this satisfy causal order?

No

Example



Does this satisfy FIFO order?

No


Recap: Ordered Multicast

- FIFO ordering
 - If a correct process issues $\text{multicast}(g, m)$ and then $\text{multicast}(g, m')$, then every correct process that delivers m' will have already delivered m .
- Causal ordering
 - If $\text{multicast}(g, m) \rightarrow \text{multicast}(g, m')$ then any correct process that delivers m' will have already delivered m .
 - Note that \rightarrow counts multicast messages delivered to the application, rather than all network messages.
- Total ordering

Total Order

- Ensures all processes deliver all multicasts in the same order.
- Unlike FIFO and causal, this does not pay attention to order of multicast sending.
- Formally
 - If a correct process delivers message m before m' (independent of sending order), then any other correct process that delivers m' will have already delivered m .

Total Order: Example



The order of receipt of multicasts is the same at all processes.

M1:1, then M2:1, then M3:1, then M3:2

May need to delay delivery of some messages.

Causal vs Total

- Total ordering does not imply causal ordering.
- Causal ordering does not imply total ordering.


Hybrid variants

- We can have hybrid ordering protocols:
 - Causal-total hybrid protocol satisfies both Causal and total orders.

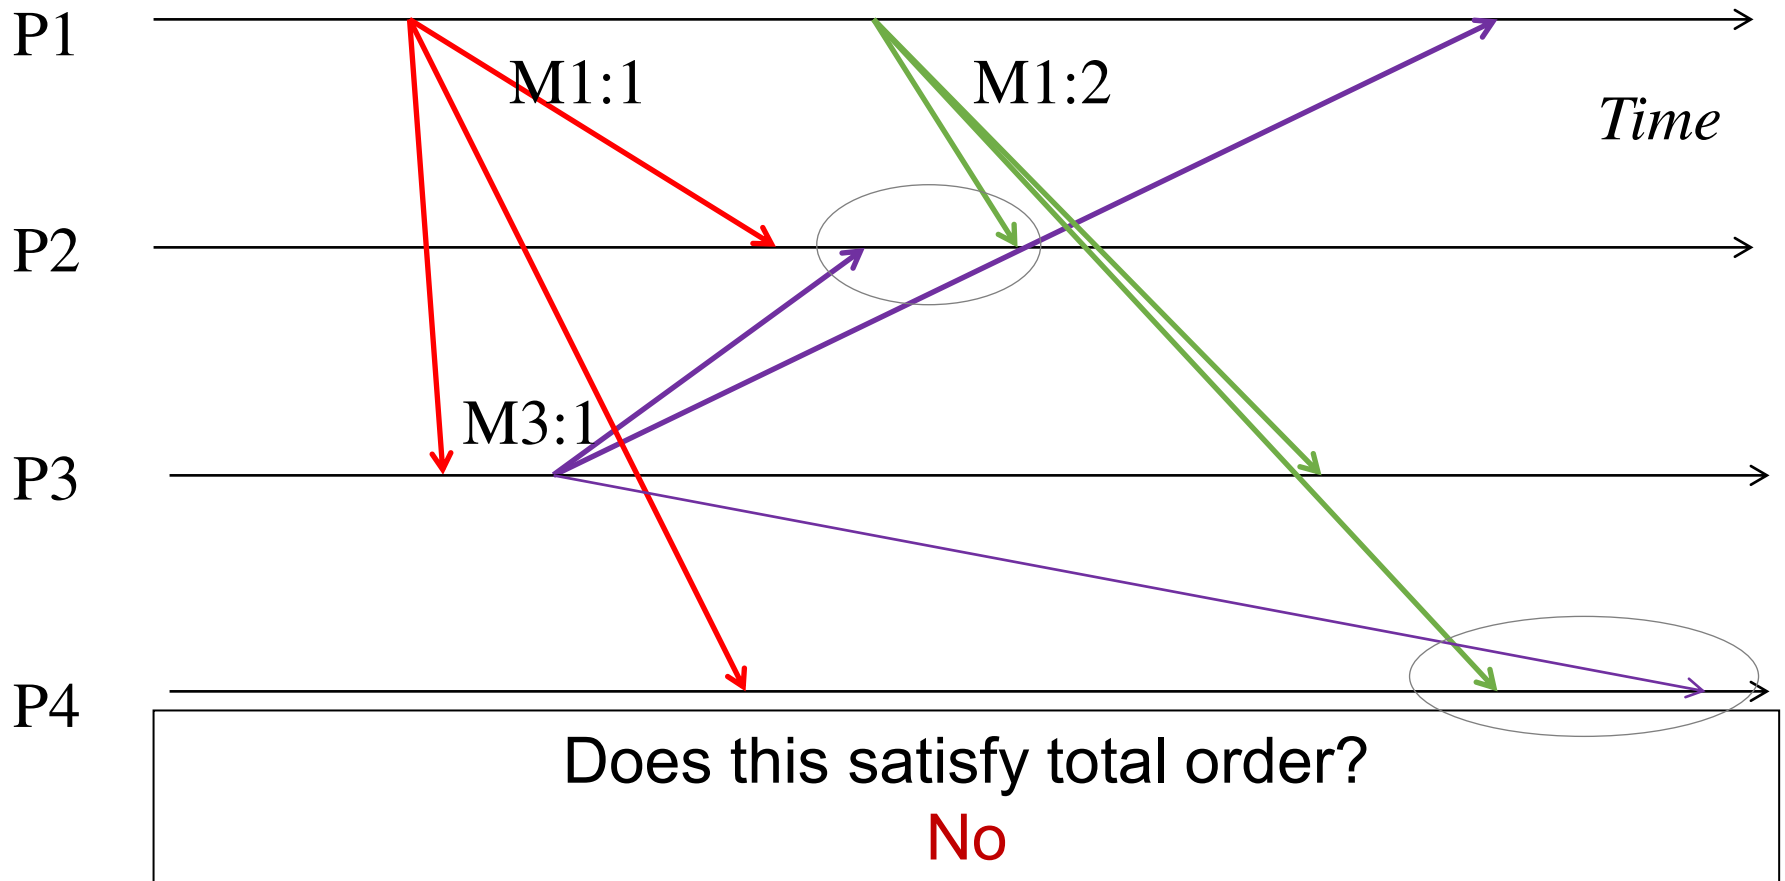
Ordered Multicast

- FIFO ordering: If a correct process issues $\text{multicast}(g, m)$ and then $\text{multicast}(g, m')$, then every correct process that delivers m' will have already delivered m .
- Causal ordering: If $\text{multicast}(g, m) \rightarrow \text{multicast}(g, m')$ then any correct process that delivers m' will have already delivered m .
 - Note that \rightarrow counts messages delivered to the application, rather than all network messages.
- Total ordering: If a correct process delivers message m before m' , then any other correct process that delivers m' will have already delivered m .


Example



Example



Example



Does this satisfy total order?

Yes

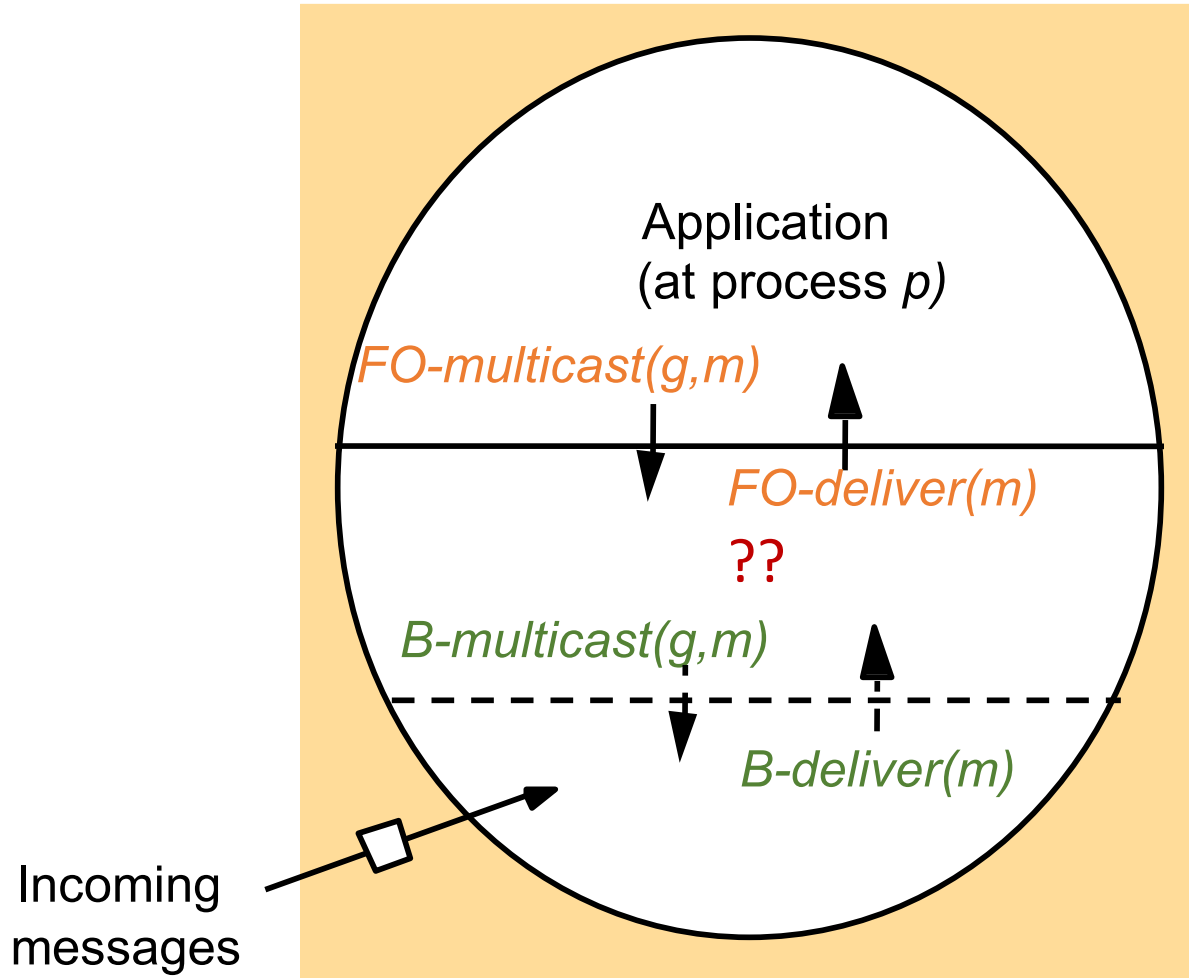
Next Question

How do we implement ordered multicast?

Ordered Multicast

- FIFO ordering
 - If a correct process issues $\text{multicast}(g, m)$ and then $\text{multicast}(g, m')$, then every correct process that delivers m' will have already delivered m .
- Causal ordering
 - If $\text{multicast}(g, m) \rightarrow \text{multicast}(g, m')$ then any correct process that delivers m' will have already delivered m .
 - Note that \rightarrow counts multicast messages delivered to the application, rather than all network messages.
- Total ordering
 - If a correct process delivers message m before m' then any other correct process that delivers m' will have already delivered m .

Implementing FIFO order multicast




Implementing FIFO order multicast

- Each receiver maintains a per-sender sequence number
 - Processes P_1 through P_N
 - P_i maintains a vector of sequence numbers $P_i[1 \dots N]$ (initially all zeroes)
 - $P_i[j]$ is latest sequence number P_i has received from P_j

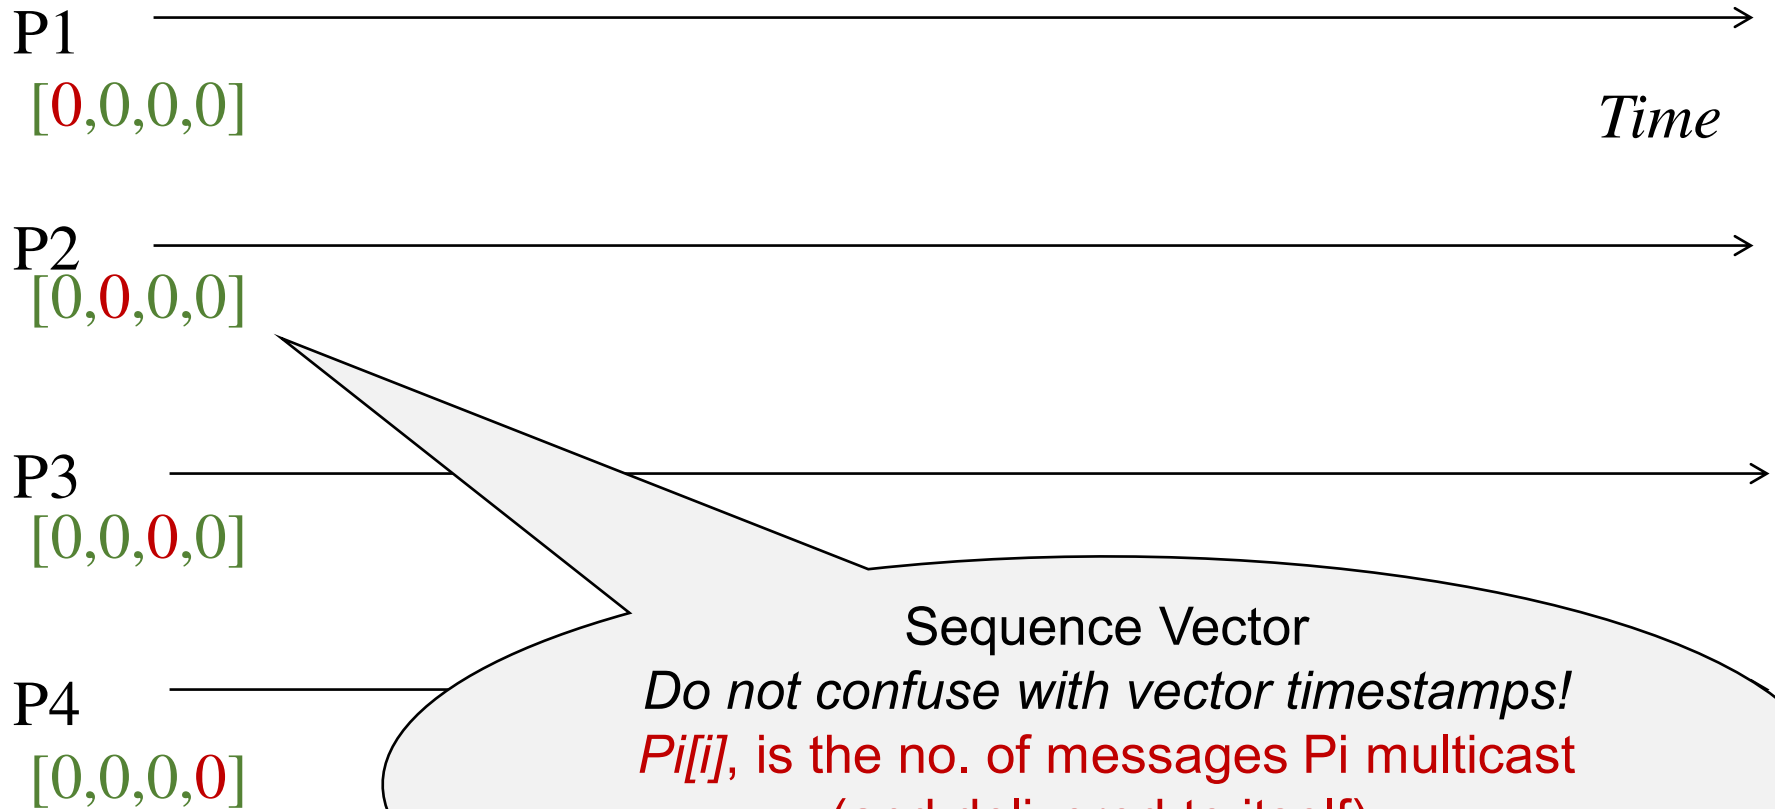
Implementing FIFO order multicast

- On FO-multicast(g, m) at process P_j :
 - set $P_j[j] = P_j[j] + 1$
 - piggyback $P_j[j]$ with m as its sequence number.
 - B-multicast($g, \{m, P_j[j]\}$)
- On B-deliver($\{m, S\}$) at P_i from P_j : *If P_i receives a multicast from P_j with sequence number S in message*
 - if ($S == P_i[j] + 1$) then
 - FO-deliver(m) to application
 - set $P_i[j] = P_i[j] + 1$
 - else buffer this multicast until above condition is true

FIFO order multicast execution



FIFO order multicast execution



Sequence Vector


Do not confuse with vector timestamps!
 *$P_i[j]$, is the no. of messages P_i multicast
(and delivered to itself).*

*$P_i[j] \forall j \neq i$ is no. of messages delivered at
 P_i from P_j .*

FIFO order multicast execution




FIFO order multicast execution




Self-deliveries omitted for simplicity.


FIFO order multicast execution



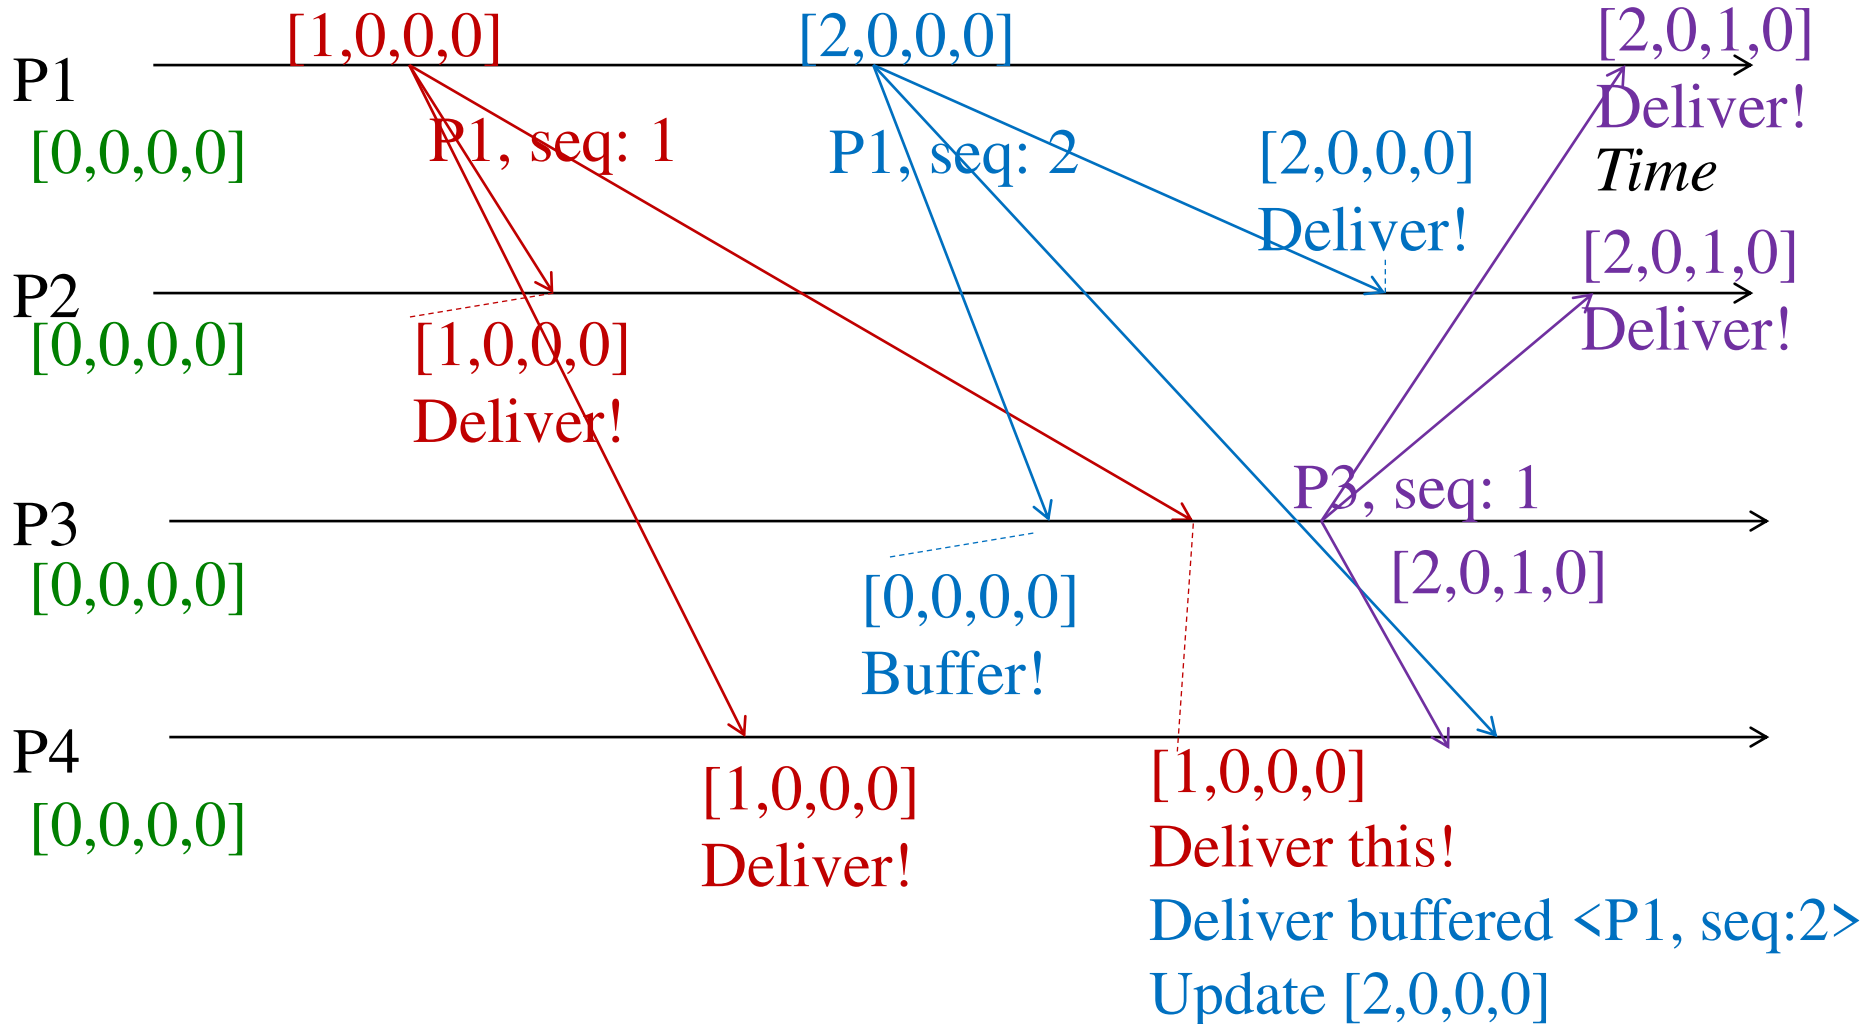
FIFO order multicast execution




FIFO order multicast execution



FIFO order multicast execution



FIFO order multicast execution



Implementing FIFO order multicast

- On FO-multicast(g, m) at process P_j :
 - set $P_j[j] = P_j[j] + 1$
 - piggyback $P_j[j]$ with m as its sequence number.
 - B-multicast($g, \{m, P_j[j]\}$)
- On B-deliver($\{m, S\}$) at P_i from P_j : *If P_i receives a multicast from P_j with sequence number S in message*
 - if ($S == P_i[j] + 1$) then
 - FO-deliver(m) to application
 - set $P_i[j] = P_i[j] + 1$
 - else buffer this multicast until above condition is true

Implementing FIFO reliable multicast

- On FO-multicast(g, m) at process P_j :
 - set $P_j[j] = P_j[j] + 1$
 - piggyback $P_j[j]$ with m as its sequence number.
 - R-multicast($g, \{m, P_j[j]\}$)
- On R-deliver($\{m, S\}$) at P_i from P_j : *If P_i receives a multicast from P_j with sequence number S in message*
 - if ($S == P_i[j] + 1$) then
 - FO-deliver(m) to application
 - set $P_i[j] = P_i[j] + 1$
 - else buffer this multicast until above condition is true

Ordered Multicast

- FIFO ordering: If a correct process issues $\text{multicast}(g, m)$ and then $\text{multicast}(g, m')$, then every correct process that delivers m' will have already delivered m .
- Causal ordering: If $\text{multicast}(g, m) \rightarrow \text{multicast}(g, m')$ then any correct process that delivers m' will have already delivered m .
 - Note that \rightarrow counts multicast messages delivered to the application, rather than all network messages.
- Total ordering: If a correct process delivers message m before m' then any other correct process that delivers m' will have already delivered m .

Implementing total order multicast

- Basic idea:
 - Same sequence number counter across different processes.
 - Instead of different sequence number counter for each process.
- Two types of approach
 - Using a centralized sequencer
 - A decentralized mechanism (ISIS)

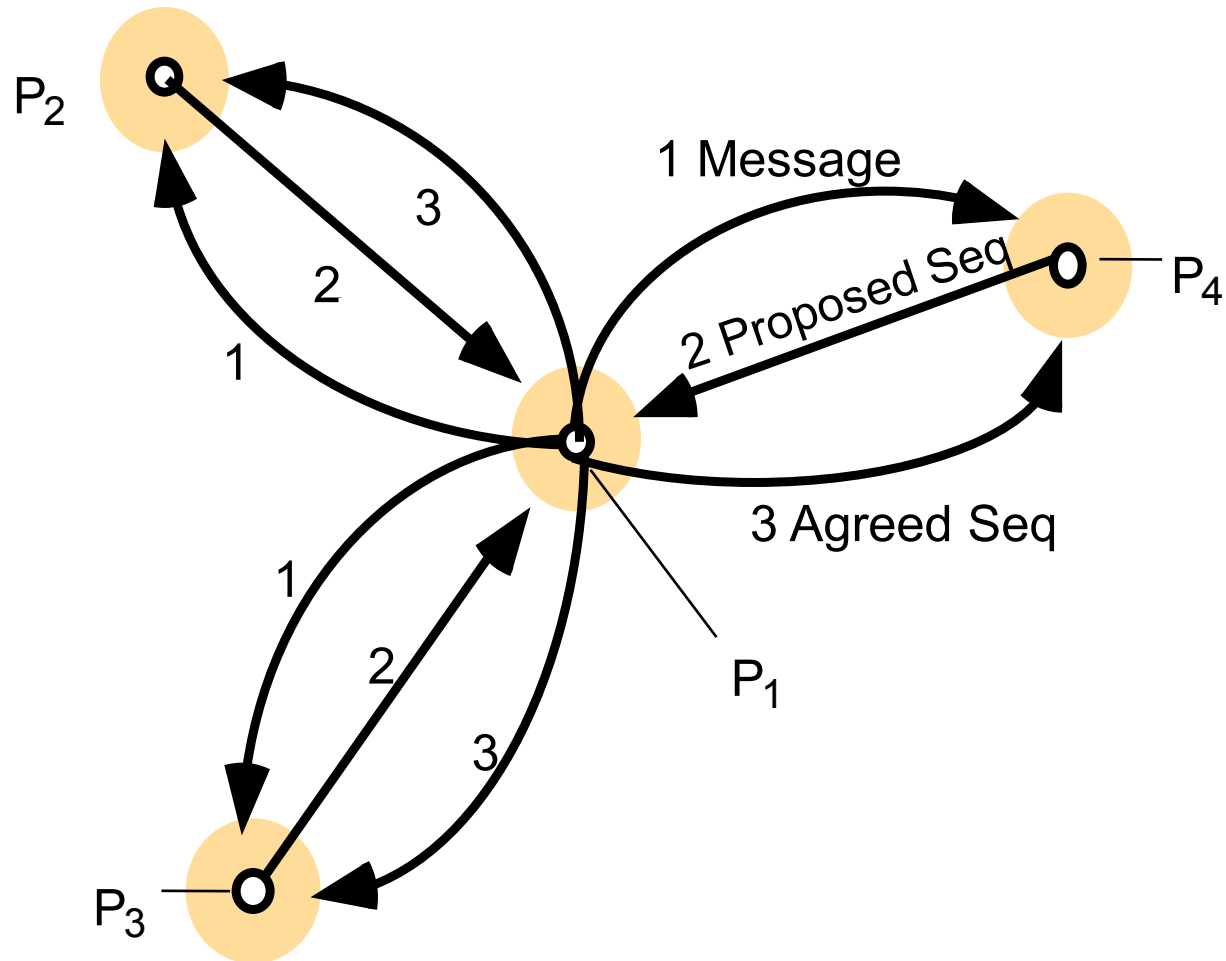
Sequencer based total ordering

- Special process elected as leader or sequencer.
- TO-multicast(g, m) at P_i :
 - B-multicast message m to group g *and the sequencer*
- Sequencer:
 - Maintains a global sequence number S (initially 0)
 - When a multicast message m is B-delivered to it:
 - sets $S = S + 1$, and B-multicast($g, \{\text{"order"}, m, S\}$)
- Receive multicast at process P_i :
 - P_i maintains a local received global sequence number S_i (initially 0)
 - On B-deliver(m) at P_i from P_j , buffers it until both conditions satisfied
 1. B-deliver($\{\text{"order"}, m, S\}$) at P_i from sequencer, and
 2. $S_i + 1 = S$
 - Then TO-deliver(m) to application and set $S_i = S_i + 1$

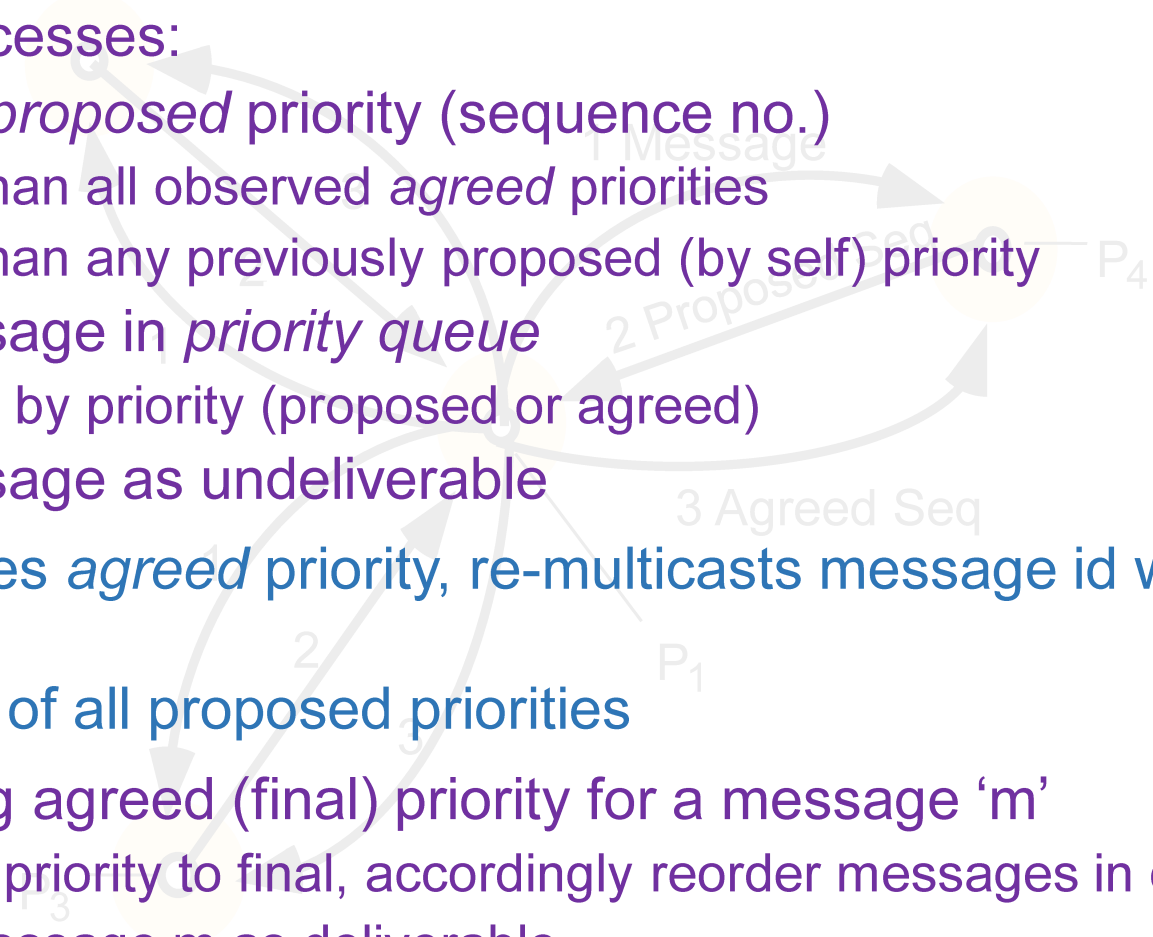
Implementing total order multicast

- Basic idea:
 - Same sequence number counter across different processes.
 - Instead of different sequence number counter for each process.
- Two types of approach
 - Using a centralized sequencer
 - A decentralized mechanism (ISIS)


ISIS algorithm for total ordering



ISIS algorithm for total ordering

- Sender multicasts message to everyone.
 - Receiving processes:
 - reply with *proposed* priority (sequence no.)
 - larger than all observed *agreed* priorities
 - larger than any previously proposed (by self) priority
 - store message in *priority queue*
 - ordered by priority (proposed or agreed)
 - mark message as undeliverable
 - Sender chooses *agreed* priority, re-multicasts message id with agreed priority
 - maximum of all proposed priorities
 - Upon receiving agreed (final) priority for a message 'm'
 - Update m's priority to final, accordingly reorder messages in queue.
 - mark the message m as deliverable.
 - deliver any deliverable messages at front of priority queue.
- 
- The diagram illustrates the ISIS algorithm for total ordering across four processes: P1, P2, P3, and P4. It shows the flow of messages and the evolution of priorities. 1. Initial State: P1 has a message 'm' with priority 1. 2. Multicast: P1 multicasts the message to P2, P3, and P4. 3. Proposals: P2 proposes priority 2, P3 proposes priority 3, and P4 proposes priority 4. 4. Agreed Priority: The sender (P1) chooses the maximum of all proposed priorities (4) as the agreed priority. 5. Re-multicast: P1 re-multicasts the message with the agreed priority (4). 6. Final State: All processes (P1, P2, P3, P4) have the message 'm' with the agreed priority (4) and mark it as deliverable. The priority queue is ordered by priority, and deliverable messages are delivered at the front.


Example: ISIS algorithm



How do we break ties?

- Problem: priority queue requires unique priorities.
- Solution: add process # to suggested priority.
 - $\text{priority.}(\textit{id of the process that proposed the priority})$
 - i.e., $3.2 == \text{process 2 proposed priority 3}$
- Compare on priority first, use process # to break ties.
 - $2.1 > 1.3$
 - $3.2 > 3.1$

Example: ISIS algorithm



Ordered Multicast

- FIFO ordering
 - If a correct process issues $\text{multicast}(g, m)$ and then $\text{multicast}(g, m')$, then every correct process that delivers m' will have already delivered m .
- Causal ordering
 - If $\text{multicast}(g, m) \rightarrow \text{multicast}(g, m')$ then any correct process that delivers m' will have already delivered m .
 - Note that \rightarrow counts multicast messages delivered to the application, rather than all network messages.
- Total ordering
 - If a correct process delivers message m before m' , then any other correct process that delivers m' will have already delivered m .

To be continued in next class

- Proof of total-ordering with ISIS.
- Implementation of causal order multicast.

Summary

- Multicast is an important communication mode in distributed systems.
- Applications may have different requirements:
 - Reliability
 - Ordering: FIFO, Causal, Total
 - Combinations of the above.