

Distributed Systems

ECE428

Lecture 9

Adopted from Spring 2021

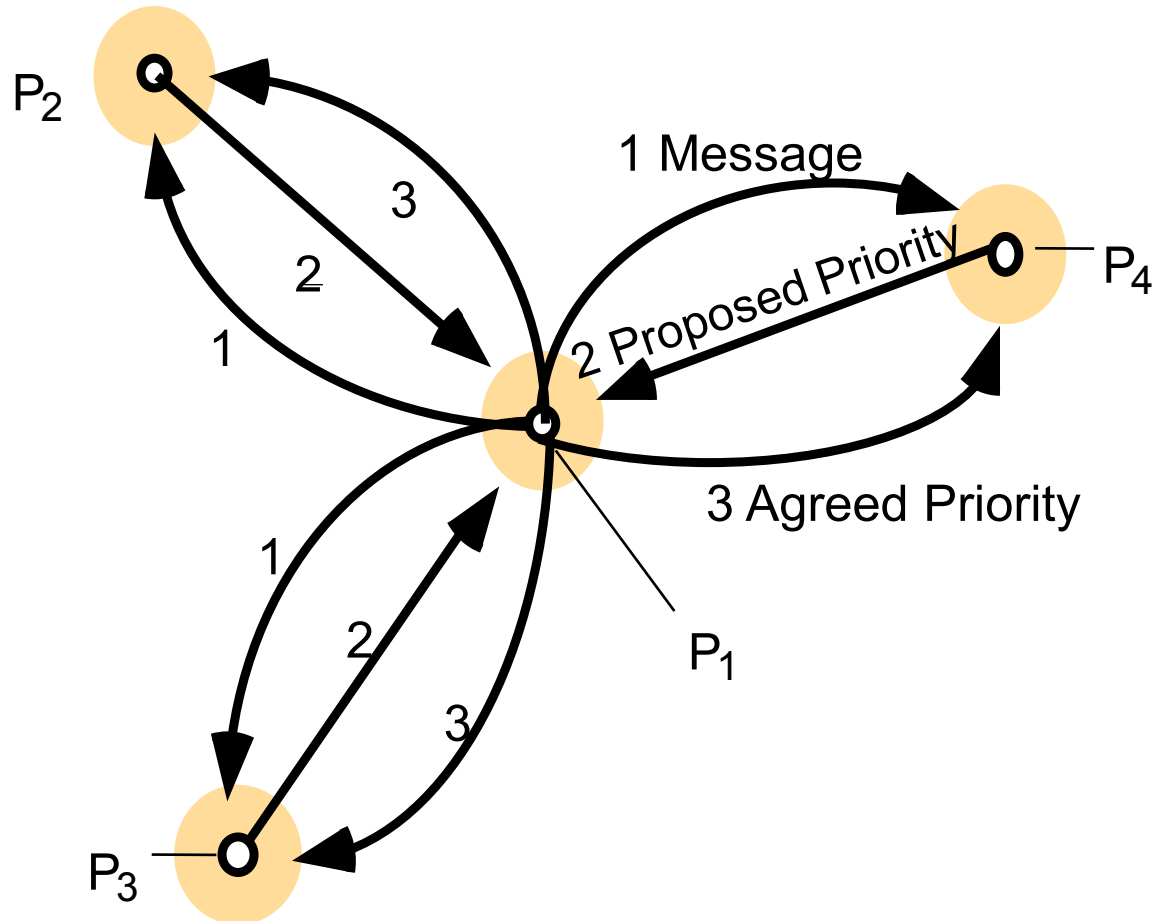
Today's agenda

- Wrap up Multicast
 - Chapter 15.4
 - Tree-based multicast and Gossip
- Mutual Exclusion
 - Chapter 15.2

Recap: Ordered Multicast

- FIFO ordering: If correct process issues $\text{multicast}(g, m)$ and then $\text{multicast}(g, m')$, then every correct process that delivers m' will have already delivered m .
- Causal ordering: If $\text{multicast}(g, m) \rightarrow \text{multicast}(g, m')$ then any correct process that delivers m' will have already delivered m .
 - Note that \rightarrow counts multicast messages delivered to the application, rather than all network messages.
- Total ordering: If a correct process delivers message m before m' , then any other correct process that delivers m' will have already delivered m .

ISIS algorithm for total ordering



Proposed Priority: *higher than all priorities proposed by the process and agreed priorities received by the process so far.*

Agreed Priority: *Maximum of all proposed priority for the message*

Proof of total order with ISIS

- Consider messages, m_1 and m_2 , and two processes, p and p' .
- Suppose that p delivers m_1 before m_2 .
- When p delivers m_1 , it is at head of the queue. m_2 is either:
 - Already in p 's queue, and deliverable, so
 - $\text{Final_priority}(m_1) < \text{Final_priority}(m_2)$
 - Already in p 's queue, and not deliverable, so
 - $\text{Final_priority}(m_1) < \text{Proposed_priority}(m_2) \leq \text{Final_priority}(m_2)$
 - Not yet in p 's queue:
 - same as above, since proposed priority $>$ priority of any delivered message
- Suppose p' delivers m_2 before m_1 , by the same argument:
 - $\text{Final_priority}(m_2) < \text{Final_priority}(m_1)$
 - Contradiction!

Ordered Multicast

- FIFO ordering
 - If a correct process issues $\text{multicast}(g, m)$ and then $\text{multicast}(g, m')$, then every correct process that delivers m' will have already delivered m .
- Causal ordering
 - If $\text{multicast}(g, m) \rightarrow \text{multicast}(g, m')$ then any correct process that delivers m' will have already delivered m .
 - Note that \rightarrow counts multicast messages delivered to the application, rather than all network messages.
- Total ordering
 - If a correct process delivers message m before m' then any other correct process that delivers m' will have already delivered m .

Implementing causal order multicast

- Similar to FIFO Multicast
 - What you send with a message differs.
 - Updating rules differ.
- Each receiver maintains a vector of per-sender sequence numbers (integers)
 - Processes P_1 through P_N .
 - P_i maintains a vector of sequence numbers $P_i[1 \dots N]$ (initially all zeroes).
 - $P_i[j]$ is the latest sequence number P_i has received from P_j .


... these are NOT vector logical clocks!

Implementing causal order multicast

- *CO-multicast*(g, m) at P_j :
 - set $P_j[j] = P_j[j] + 1$
 - piggyback entire vector $P_j[1 \dots N]$ with m .
 - $B\text{-multicast}(g, \{m, P_j[1 \dots N]\})$
- On $B\text{-deliver}(\{m, V[1 \dots N]\})$ at P_i from P_j : If P_i receives a multicast from P_j with sequence vector $V[1 \dots N]$, buffer it until both conditions are true:
 1. This message is next one P_i is expecting from P_j , i.e.,
$$V[j] = P_i[j] + 1$$
 2. All multicasts, anywhere in the group, which happened before m have been received at P_i , i.e.,
$$\text{For all } k \neq j: V[k] \leq P_i[k]$$


When these conditions satisfied, $CO\text{-deliver}(m)$, and set $P_i[j] = V[j]$

Causal order multicast execution




Self-deliveries omitted for simplicity.

Causal order multicast execution



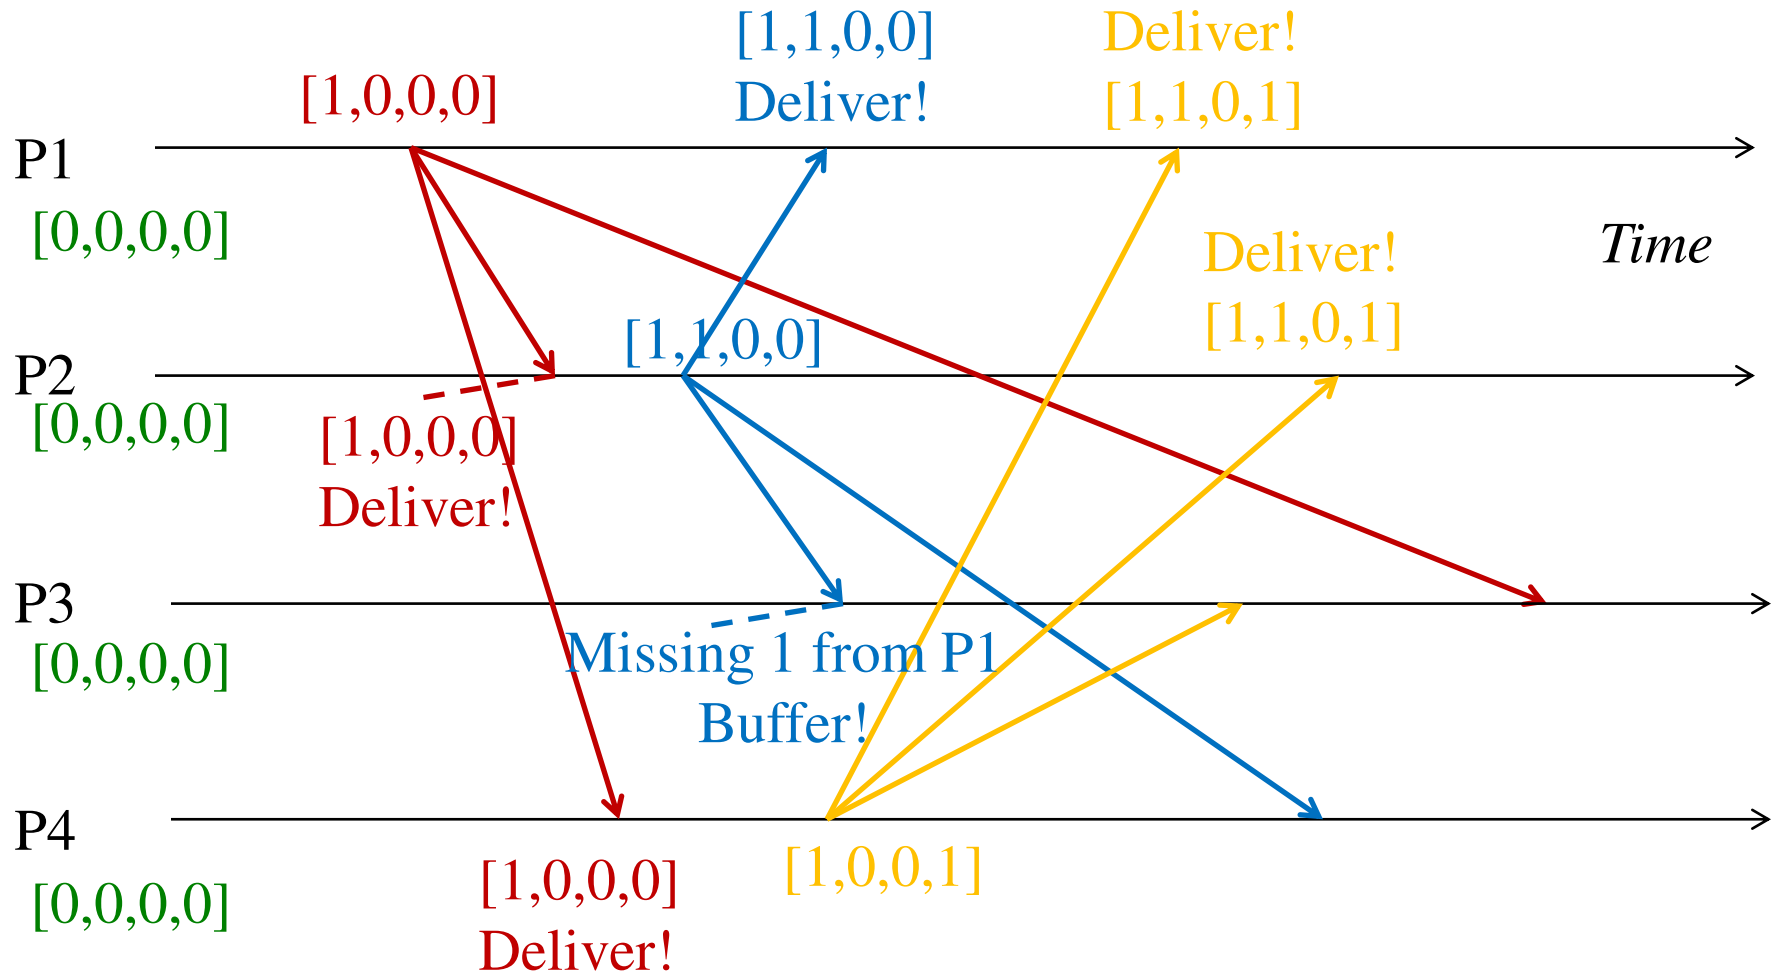
Self-deliveries omitted for simplicity.

Causal order multicast execution



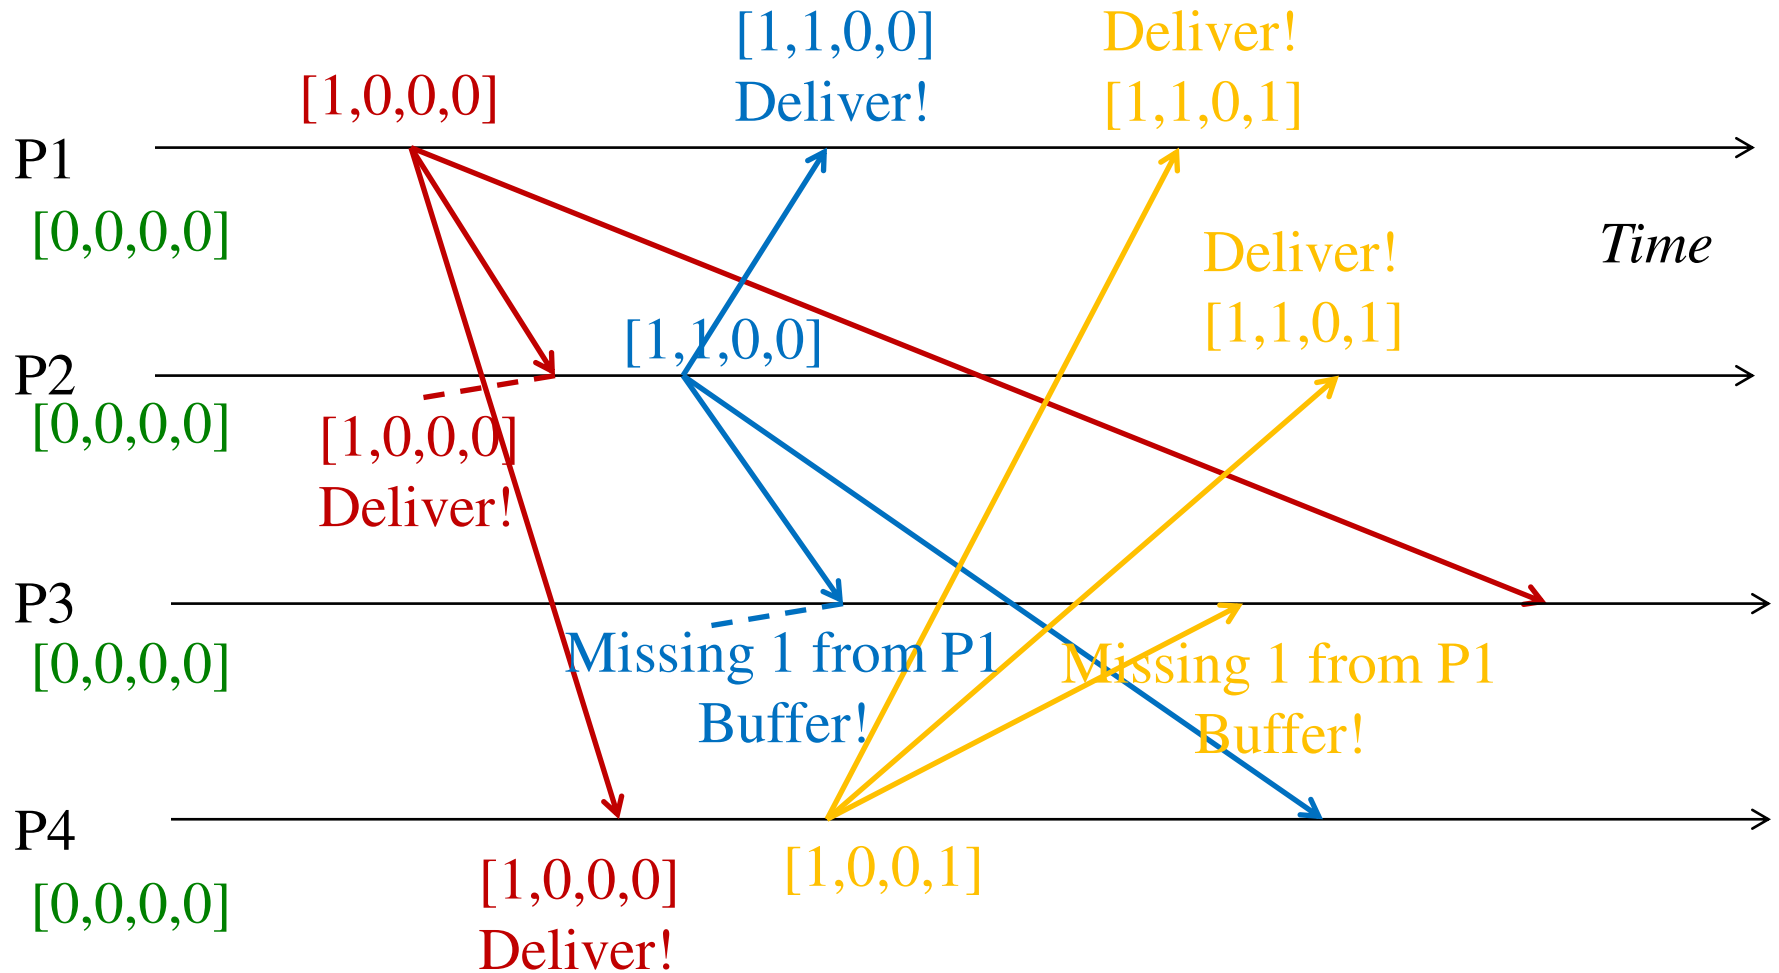
Self-deliveries omitted for simplicity.

Causal order multicast execution

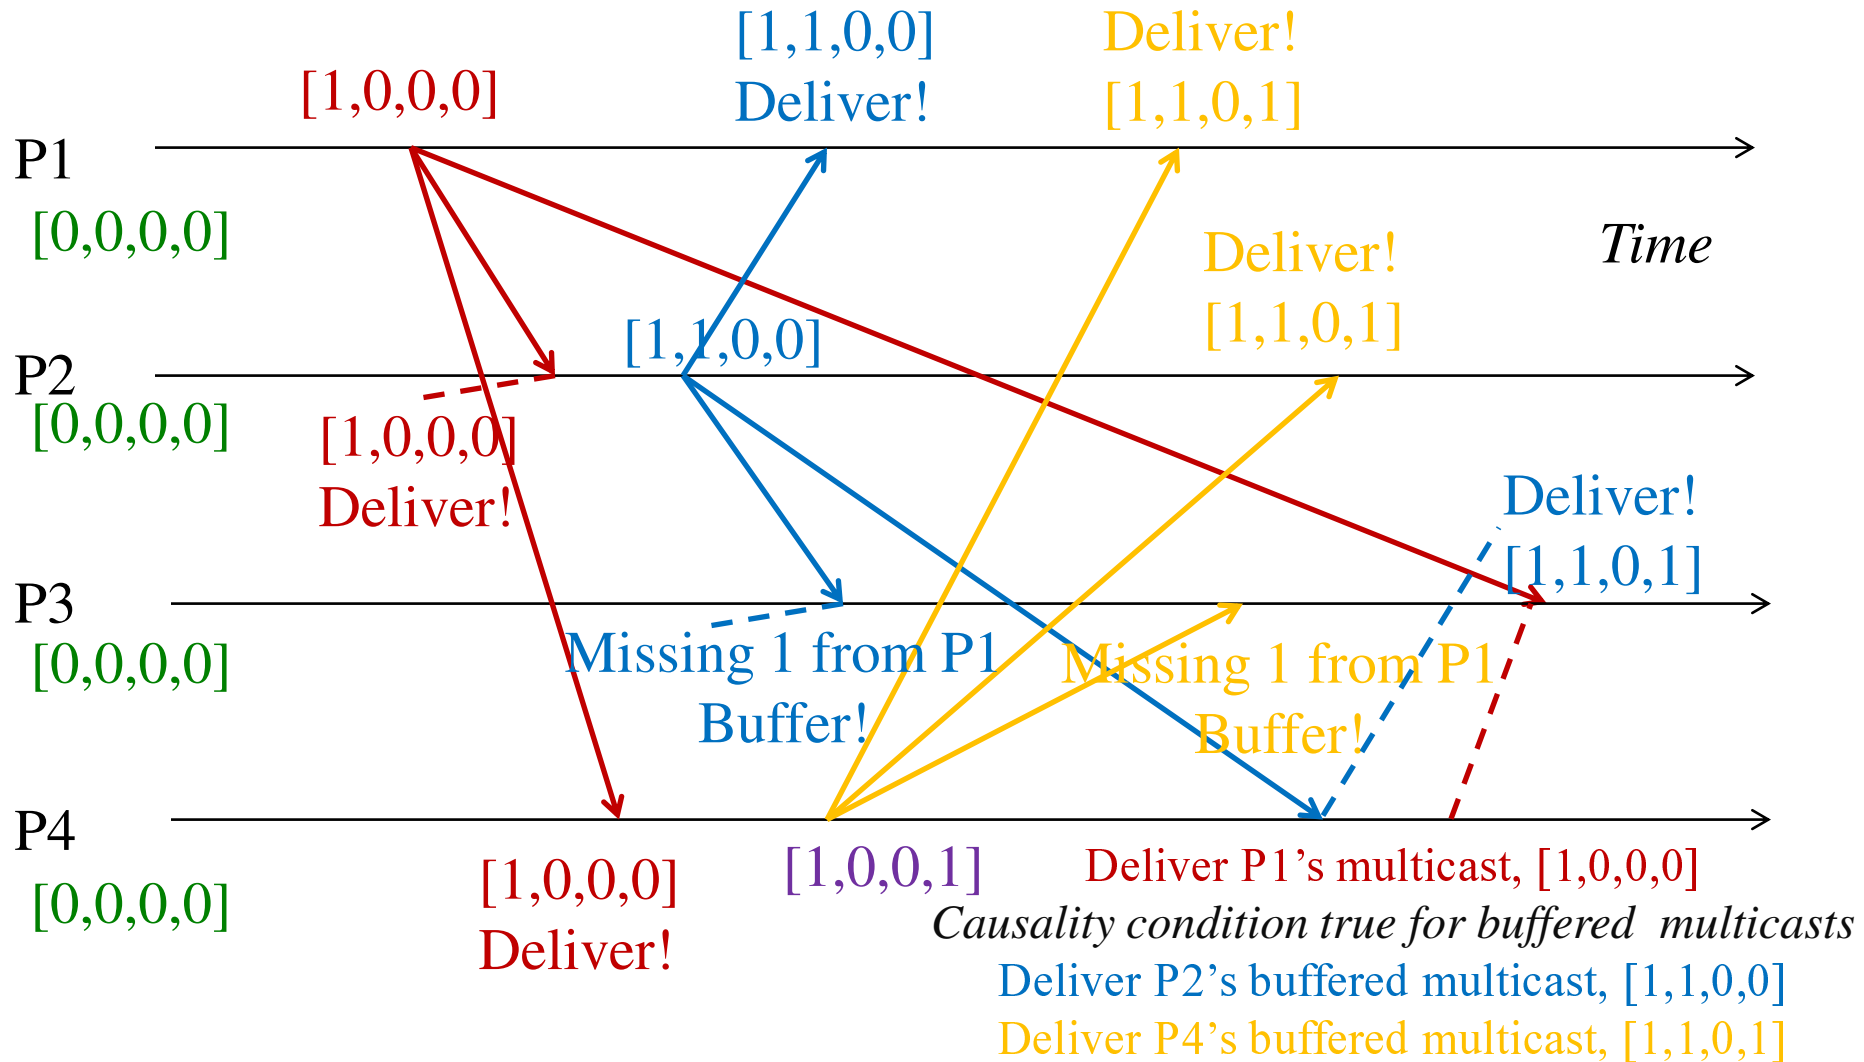


Self-deliveries omitted for simplicity.

Causal order multicast execution



Causal order multicast execution



Causal order multicast implementation

- Only looks at multicast messages delivered to the application.
- Ignores causality created due to other network messages.

Ordered Multicast






- FIFO ordering
 - If a correct process issues $\text{multicast}(g, m)$ and then $\text{multicast}(g, m')$, then every correct process that delivers m' will have already delivered m .
- Causal ordering
 - If $\text{multicast}(g, m) \rightarrow \text{multicast}(g, m')$ then any correct process that delivers m' will have already delivered m .
 - Note that \rightarrow counts multicast messages delivered to the application, rather than all network messages.
- Total ordering
 - If a correct process delivers message m before m' , then any other correct process that delivers m' will have already delivered m .

More efficient multicast mechanisms


- Our focus so far has been on the application-level semantics of multicast.
- *What are some of the more efficient underlying mechanisms for a B-multicast?*

B-Multicast

Sender




B-Multicast using unicast sends




B-Multicast using unicast sends

Sender

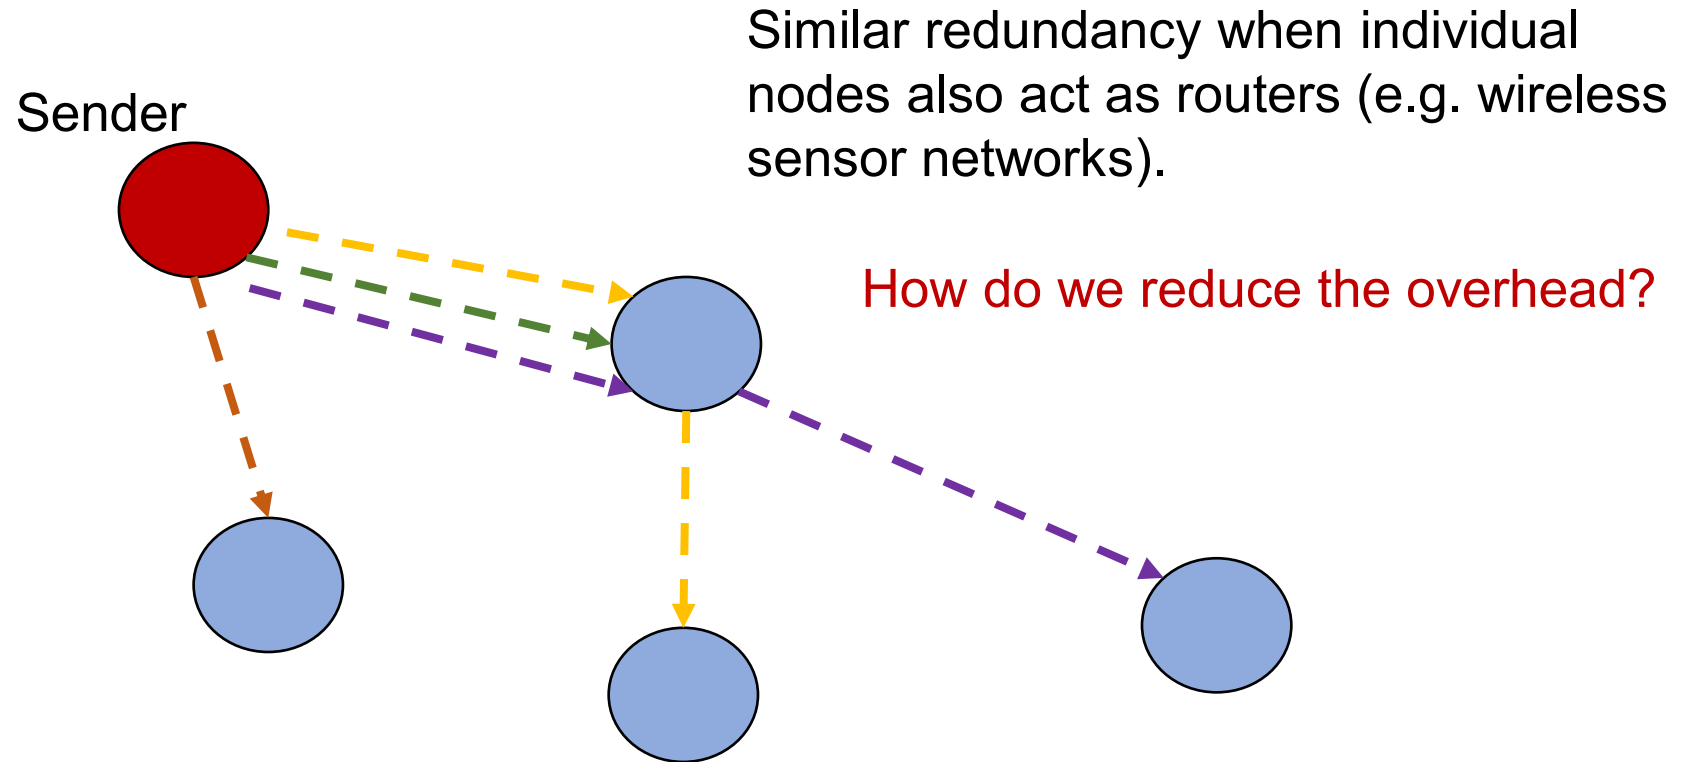
Closer look at physical network paths.



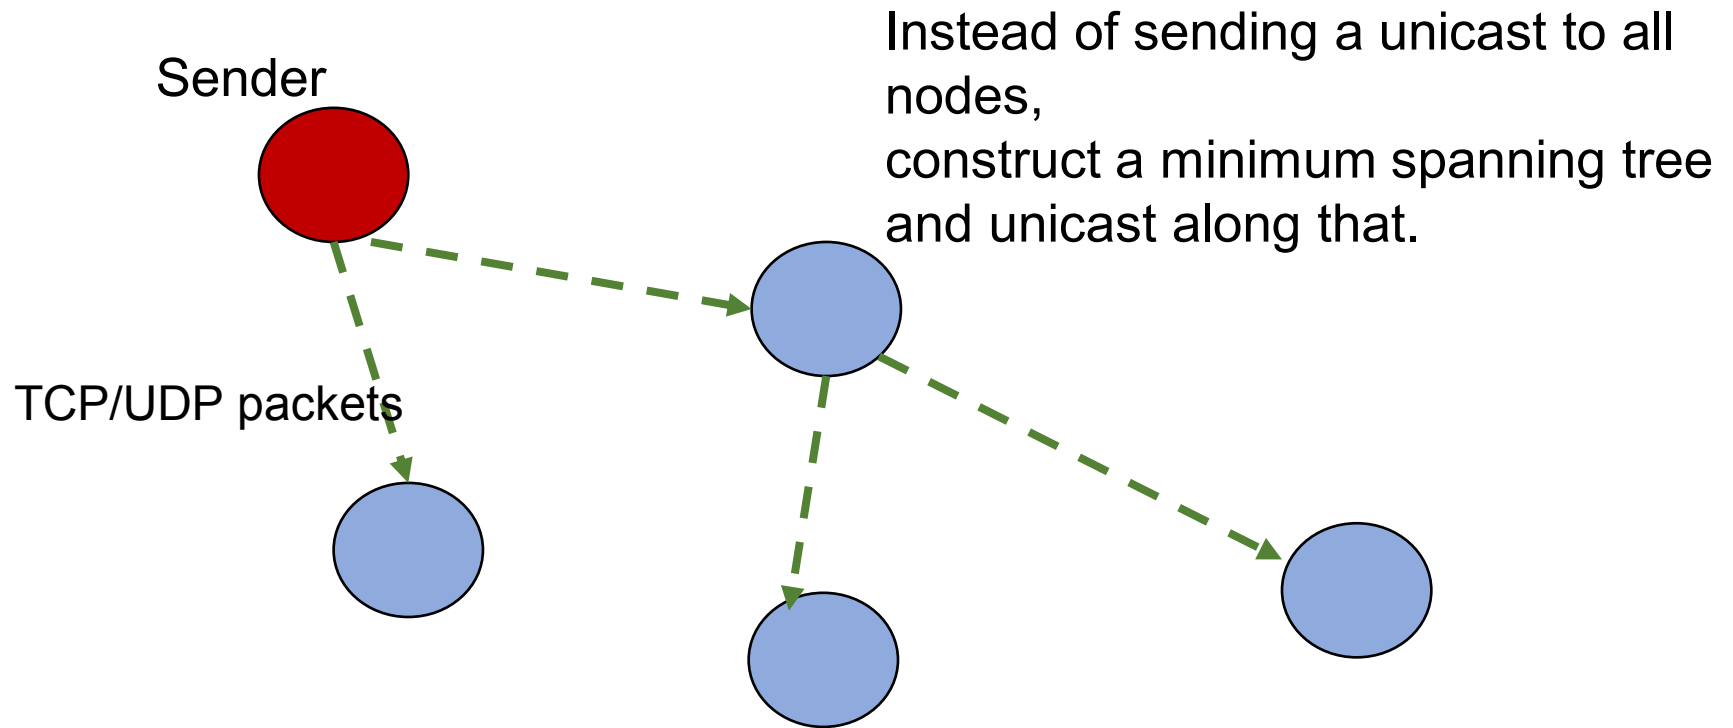
B-Multicast using unicast sends



B-Multicast using unicast sends




Tree-based multicast



Tree-based multicast

A process does not directly send messages to *all* other processes in the group.




Tree-based multicast

A process does not directly send messages to *all* other processes in the group.

Sender


It sends a message to only a subset of processes.

Closer look at the physical network.




Tree-based multicast

Sender




Also possible to construct a tree that includes network routers. IP multicast!

Tree-based multicast




Third approach: Gossip

Transmit to b random targets.




Third approach: Gossip

Transmit to b random targets.



Third approach: Gossip

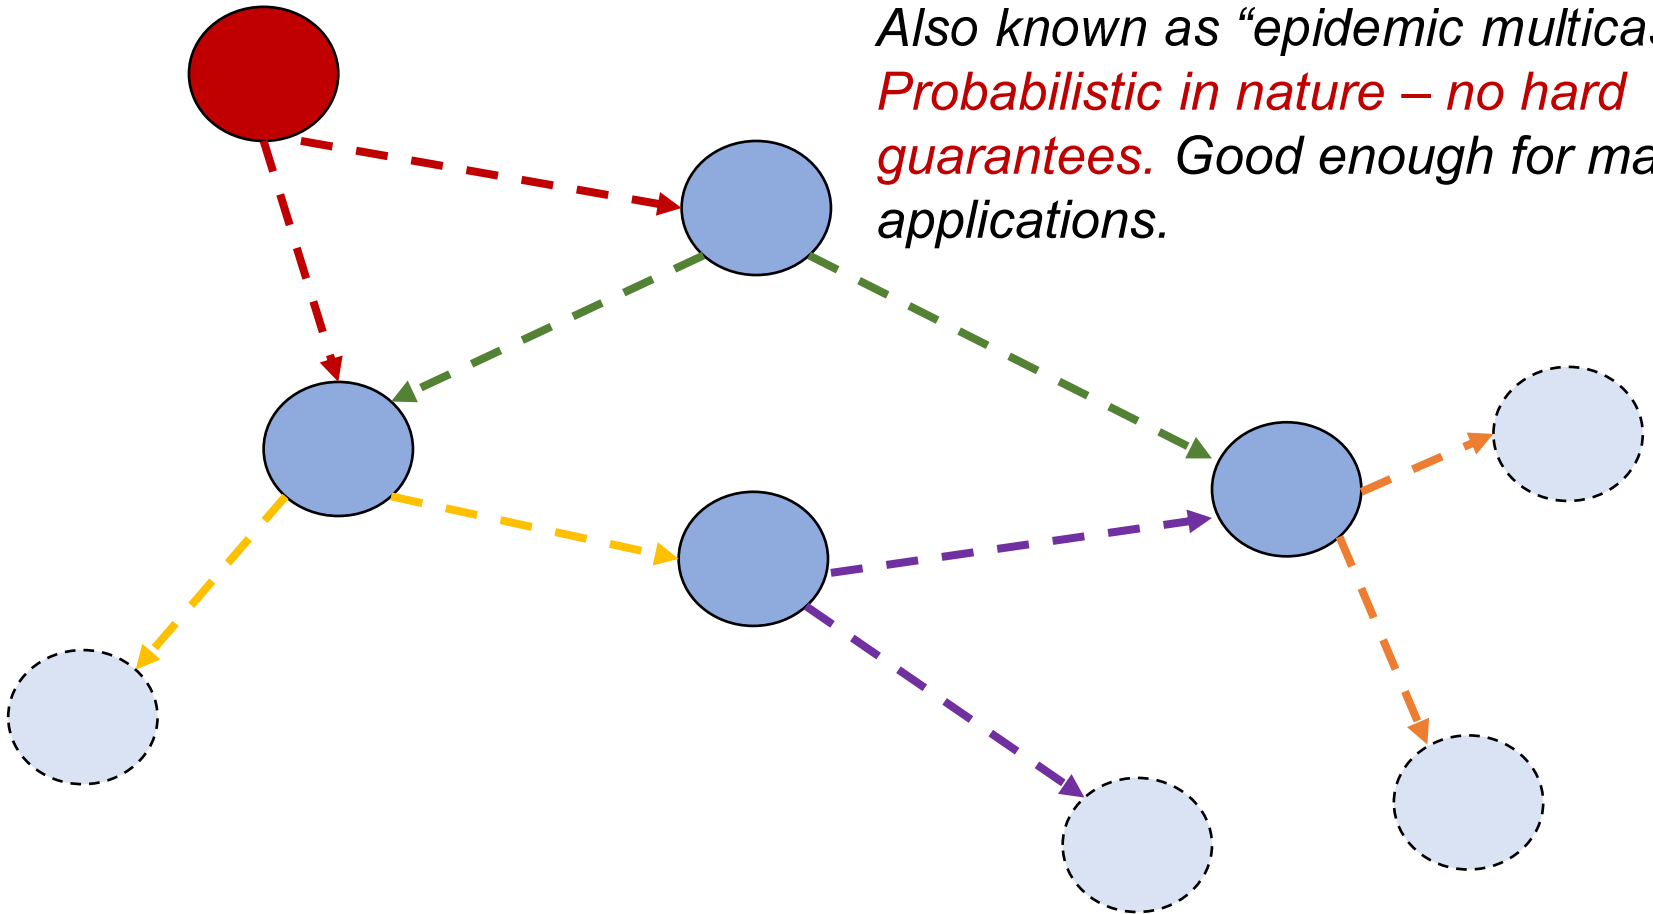
Transmit to b random targets.



Third approach: Gossip

No “tree-construction” overhead.
More efficient than unicasting to all
receivers.


*Also known as “epidemic multicast”.
**Probabilistic in nature – no hard
guarantees.** Good enough for many
applications.*



Third approach: Gossip

Used in many real-world systems:

- Facebook's distributed datastore uses it to determine group membership and failures.
- Bitcoin uses it to exchange transaction information between nodes.



Multicast Summary

- Multicast is important for applications in distributed systems.
- Applications may have different requirements:
 - Basic
 - Reliable
 - Ordering: FIFO, Causal, Total
 - Combinations of the above.
- Underlying mechanisms to spread the information:
 - Unicast to all receivers.
 - Tree-based multicast, and gossip: sender unicasts messages to only a subset of other processes, and they spread the message further.
 - Gossip is more scalable and more robust to failures.

Today's agenda

- Wrap up Multicast
 - Chapter 15.4
 - Tree-based multicast and Gossip
- Mutual Exclusion
 - Chapter 15.2
- Goal: reason about ways in which different processes in a distributed system can safely manipulate shared resources.

Why Mutual Exclusion?

- **Bank's Servers in the Cloud:** Two of your customers make simultaneous deposits of \$1,000 into your bank account, each from a separate ATM.
 - Both ATMs read initial amount of \$10,000 in your account concurrently from the bank's cloud server
 - Both ATMs add \$1,000 to this amount (locally at ATM)
 - Both write the final amount to the server
 - **What's wrong?**

Why Mutual Exclusion?

- **Bank's Servers in the Cloud:** Two of your customers make simultaneous deposits of \$1,000 into your bank account, each from a separate ATM.
 - Both ATMs read initial amount of \$10,000 in your account concurrently from the bank's cloud server
 - Both ATMs add \$1,000 to this amount (locally at ATM)
 - Both write the final amount to the server
 - **You lost \$1,000!**
- **The ATMs need *mutually exclusive* access to your account entry at the server**
 - or, mutually exclusive access to executing the code that modifies the account entry.

More uses of mutual exclusion

- Distributed file systems
 - Locking of files and directories
- Accessing objects in a safe and consistent way
 - Ensure at most one server has access to object at any point of time
- In industry
 - Chubby is Google's locking service

Problem statement for mutual excls.

- *Critical Section Problem*:
 - Piece of code (at all processes) for which we need to ensure there is at most one process executing it at any point of time.
- Each process calls three functions
 - **enter()** in order to enter the critical section (CS)
 - **AccessResource()** to run the critical section code
 - **exit()** to exit the critical section

Our bank example

ATM1:

```
enter();  
    // AccessResource()  
obtain bank amount;  
add in deposit;  
update bank amount;  
    // AccessResource() end  
exit();
```

ATM2:

```
enter();  
    // AccessResource()  
obtain bank amount;  
add in deposit;  
update bank amount;  
    // AccessResource() end  
exit();
```

Mutual exclusion for a single OS

- If all processes are running in one OS on the same machine (or VM):
 - Semaphores
 - Mutexes
 - Condition variables
 - Monitors
 - ...

Processes sharing an OS: Semaphores

- Semaphore == an integer that can only be accessed via two special functions
- Semaphore S=1; // Max number of allowed accessors.

wait(S) (or P(S) or down(S)):

while(1) { // each execution of the while loop is atomic

if (S > 0) {

S--;

break;

}

}

signal(S) (or V(S) or up(s)):

S++; // atomic

exit()

Atomic operations are supported via hardware instructions such as compare-and-swap, test-and-set, etc.

Our bank example

ATM1:

```
enter();  
    // AccessResource()  
obtain bank amount;  
add in deposit;  
update bank amount;  
    // AccessResource() end  
exit();
```

ATM2:

```
enter();  
    // AccessResource()  
obtain bank amount;  
add in deposit;  
update bank amount;  
    // AccessResource() end  
exit();
```

Our bank example

Semaphore S=1; // shared

ATM1:

```
wait(S); //enter
    // AccessResource()
obtain bank amount;
add in deposit;
update bank amount;
    // AccessResource() end
signal(S); // exit
```

ATM2:

```
wait(S); //enter
    // AccessResource()
obtain bank amount;
add in deposit;
update bank amount;
    // AccessResource() end
signal(S); // exit
```

Mutual exclusion in distributed systems

- Processes communicating by passing messages.
- Cannot share variables like semaphores!
- *How do we support mutual exclusion in a distributed system?*

Mutual exclusion in distributed systems

- Our focus today: Classical algorithms for mutual exclusion in distributed systems.
 - Central server algorithm
 - Ring-based algorithm
 - Ricart-Agrawala Algorithm
 - Maekawa Algorithm

Mutual Exclusion Requirements

- Need to guarantee 3 properties:
 - **Safety** (essential):
 - At most one process executes in CS (Critical Section) at any time.
 - **Liveness** (essential):
 - Every request for a CS is granted eventually.
 - **Ordering** (desirable):
 - Requests are granted in the order they were made.

System Model

- Each pair of processes is connected by reliable channels (such as TCP).
- Messages sent on a channel are eventually delivered to recipient, and in FIFO order.
- Processes do not fail.
 - Fault-tolerant variants exist in literature.

Mutual exclusion in distributed systems

- Our focus today: Classical algorithms for mutual exclusion in distributed systems.
 - Central server algorithm
 - Ring-based algorithm
 - Ricart-Agrawala Algorithm
 - Maekawa Algorithm

Central Server Algorithm

- Elect a central server (or leader)
- Leader keeps
 - A queue of waiting requests from processes who wish to access the CS
 - A special token which allows its holder to access CS
- Actions of any process in group:
 - **enter()**
 - Send a request to leader
 - Wait for token from leader
 - **exit()**
 - Send back token to leader

Central Server Algorithm

- Leader Actions:
 - On receiving a request from process P_i
 - if (leader has token)
 - Send token to P_i
 - else
 - Add P_i to queue
 - On receiving a token from process P_i
 - if (queue is not empty)
 - Dequeue head of queue (say P_j), send that process the token
 - else
 - Retain token

Analysis of Central Algorithm

- Safety – at most one process in CS
 - Exactly one token
- Liveness – every request for CS granted eventually
 - With N processes in system, queue has at most N processes
 - If each process exits CS eventually and no failures, liveness guaranteed
- Ordering:
 - FIFO ordering guaranteed in order of requests received at leader
 - Not in the order in which requests were sent or the order in which processes enter CS!

Analysis of Central Algorithm

- Safety – at most one process in CS
 - Exactly one token
- Liveness – every request for CS granted eventually
 - With N processes in system, queue has at most N processes
 - If each process exits CS eventually and no failures, liveness guaranteed
- Ordering:
 - FIFO ordering guaranteed in order of requests received at leader
 - Not in the order in which requests were sent or the order in which processes enter CS!

To be continued in next class

- Metrics for analyzing performance of mutual exclusion algorithms.
- Other algorithms for mutual exclusion in distributed systems.
 - Central server algorithm
 - Ring-based algorithm
 - Ricart-Agrawala Algorithm
 - Maekawa Algorithm