

Distributed Systems

ECE428

Lecture 13

Adopted from Spring 2021

While we wait....

- A process initiates Bully algorithm after detecting the leader's failure.
- What is the worst-case turn-around time?
 - Assuming no other node fails.
 - Assume timeout is computed using the knowledge of one-way message latency (T)

Today's agenda

- Wrap up leader election
 - Chapter 15.3
- Consensus

Recap: Leader Election

- In a group of processes, elect a *Leader* to undertake special tasks
 - *Let everyone know* in the group about this Leader.
- Safety condition:
 - During the run of an election, a correct process has either not yet elected a leader, or has elected process with best attributes.
- Liveness condition:
 - Election run terminates and each process eventually elects someone.
- Two classical algorithms:
 - Ring-based algorithm
 - Bully algorithm
- Difficulty of ensure both safety and liveness in an asynchronous system under failures.

Bully Algorithm

- When a process wants to initiate an election
 - if it knows its id is the highest
 - it elects itself as coordinator, then sends a *Coordinator* message to all processes with lower id's. Election is completed.
 - else
 - it initiates an election by sending an *Election* message
 - (contd.)

Bully Algorithm (2)

- **else** it initiates an election by sending an *Election* message
 - Sends it to only processes that have a *higher id than itself*.
 - **if** receives no answer within *timeout*, calls itself leader and sends *Coordinator* message to all lower id processes. Election completed.
 - **if** an answer is received, then there is some non-faulty higher process => so, wait for coordinator message. If none received after another *timeout*, start a new election run.
- A process that receives an *Election* message replies with *disagree* message, and starts its own leader election protocol (unless it has already done so)

Timeout values


- Assume the one-way message transmission time (T) is known.
- First timeout value (when the process that has initiated election waits for the first response)
 - Must be set as accurately as possible.
 - If it is too small, a lower id process can declare itself to be the coordinator even when a higher id process is alive.
 - What should be the first timeout value be, given the above assumption?
 - $2T + (\text{processing time}) \approx 2T$
- When the second timeout happens (after 'disagree' message), election is re-started.
 - A very small value will lead to extra "Election" messages.
 - A suitable option is to use the worst-case turnaround time.

Performance Analysis

- Best-case
 - Second-highest id detects leader failure
 - Highest remaining id initiates election.
 - Sends $(N-2)$ Coordinator messages
 - Turnaround time: 1 message transmission time (T)
- Worst-case: For simplicity, assume no failures after a process calls for election.
 - if any lower id process detects failure and starts election.
 - Turnaround time: 4 message transmission times ($4T$)

Bully Algorithm: Example

P2 initiates election after detecting P5's failure.



Analysis

- Best-case
 - Second-highest id detects leader failure
 - Highest remaining id initiates election.
 - Sends (N-2) Coordinator messages
 - Turnaround time: 1 message transmission time
- Worst-case: For simplicity, let no failures after process calls for election.
 - Turnaround time: 4 message transmission times
 - if any lower id process detects failure and starts election.
 - Election + (disagree & Election) + (Timeout – T) + Coordinator
 - When the process with the lowest id in the system detects failure.
 - (N-1) processes altogether begin elections, each sending messages to processes with higher ids.
 - i-th highest id process sends (i-1) election messages
 - Number of Election messages
 $= N-1 + N-2 + \dots + 1 = (N-1)*N/2 = O(N^2)$

Correctness

- In synchronous system model:
 - Set timeout accurately using known bounds on network delays and processing times.
 - Satisfies safety and liveness.
- In asynchronous system model:
 - Failure detectors cannot be both accurate and complete.
 - Either liveness and safety is violated.

Why is Election so hard?

- Because it is related to the consensus problem!
- If we could solve election, then we could solve consensus!
 - Elect a process, use its id's last bit as the consensus decision.
- But (as we will soon see) consensus is impossible in asynchronous systems, so is election!

Today's agenda

- Wrap up leader election
 - Chapter 15.3
- Consensus, goals:
 - Understand the problem of consensus
 - How to achieve consensus in synchronous system
 - Difficulty of achieving consensus in asynchronous system
 - Good-enough consensus algorithms for asynchronous systems

Agenda for the next 2 weeks

- Consensus

- Consensus in synchronous systems
 - *Chapter 15.4*
- Impossibility of consensus in asynchronous systems
 - *We will not cover the proof in details*
- Good enough consensus algorithm for asynchronous systems:
 - *Paxos made simple, Leslie Lamport, 2001*
- Other forms of consensus algorithm
 - Raft (log-based consensus)
 - Block-chains (distributed consensus)

Agenda for today

- Consensus

- Consensus in synchronous systems
 - *Chapter 15.4*
- Impossibility of consensus in asynchronous systems
 - *We will not cover the proof in details*
- A good enough consensus algorithm for asynchronous systems:
 - *Paxos made simple, Leslie Lamport, 2001*
- Other forms of consensus
 - Blockchains
 - Raft (log-based consensus)

Consensus

- Each process **proposes** a value.
- All processes must **agree** on one of the proposed values.
- Examples:
 - The generals must agree on the time of attack.
 - An object replicated across multiple servers in a distributed data store.
 - All servers must agree on the current version of the object.
 - Transaction processing on replicated servers
 - Must agree on the order in which updates are applied to an object.
 -

Consensus

- Each process **proposes** a value.
- All processes must **agree** on one of the proposed values.
- The final value can be decided based on any criteria:
 - Pick minimum of all proposed values.
 - Pick maximum of all proposed values.
 - Pick the majority (with some deterministic tie-breaking rule).
 - Pick the value proposed by the *leader*.
 - *All processes must agree on who the leader is.*
 - If reliable total-order can be achieved, pick the proposed value that gets delivered first.
 - *All process must agree on the total order.*
 -

Consensus Problem

- System of N processes (P_1, P_2, \dots, P_n)
- Each process P_i :
 - begins in an *undecided* state.
 - proposes value v_i .
 - at some point during the run of a consensus algorithm, sets a decision variable d_i and enters the *decided* state.

Required Properties

- Termination: Eventually each process sets its decision variable.
- Agreement: The decision of all correct processes is the same.
 - If P_i and P_j are correct and have entered *decided* state, then $d_i = d_j$.
- Integrity: If correct processes all proposed same value, then any correct process in decided state has chosen that value.
 - *Specific definition of integrity may vary across sources and systems.*
 - *Safeguard against algorithms that decide on a fixed constant value.*

Required Properties

- Termination: Eventually each process sets its decision variable.
- Agreement: The decision of all correct processes is the same.
 - If P_i and P_j are correct and have entered *decided* state, then $d_i = d_j$.
- Integrity: If correct processes all proposed same value, then any correct process in decided state has chosen that value.
 - *Specific definition of integrity may vary across sources and systems.*
 - *Safeguard against algorithms that decide on a fixed constant value.*

Which of these properties is liveness and which is safety?

Required Properties

- Termination: Eventually each process sets its decision variable.
 - *Liveness*
- Agreement: Decision value of all correct processes is same.
 - If P_i and P_j are correct and have entered *decided* state, then $d_i = d_j$.
 - *Safety*
- Integrity: If correct processes all proposed same value, then any correct process in the decided state has chosen that value.

How do we agree on a value?


- Ring-based leader election
 - Send proposed value along with *elected* message.
 - Turnaround time: $3NT$ worst case, $2NT$ best case (no failures).
 - T is the time taken to transmit a message on channel.
 - $O(NT \times F)$ if up to F processes fail during the election run.
 - Can we do better?
- Bully algorithm
 - Send proposed value along with the *coordinator* message.
 - Turnaround time: $4T$ in the worst case without failures.
 - More than $2FT$ if up to F processes fail during the election run.

What's the best we can do?

Consider the simplest algorithm

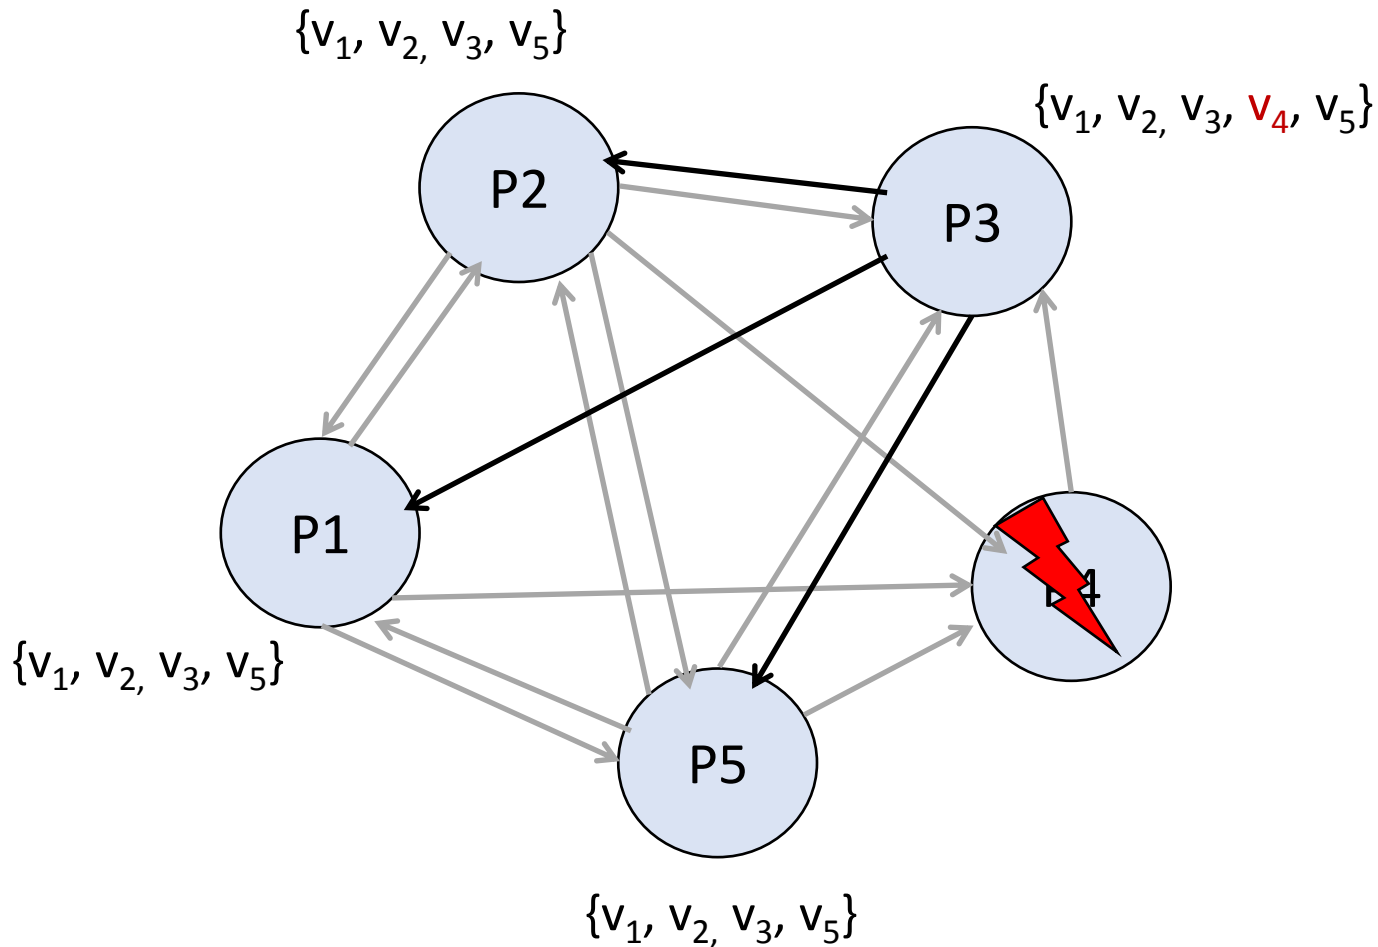
- Let's assume the system is synchronous.
- Use a simple B-multicast:
 - All processes B-multicast their proposed value to all other processes.
 - Upon receiving all proposed values, pick the minimum.
- Time taken under no failures?
 - One message transmission time (T)
- What can go wrong?
 - If we consider failures, is simple B-multicast enough?

B-multicast is not enough for this



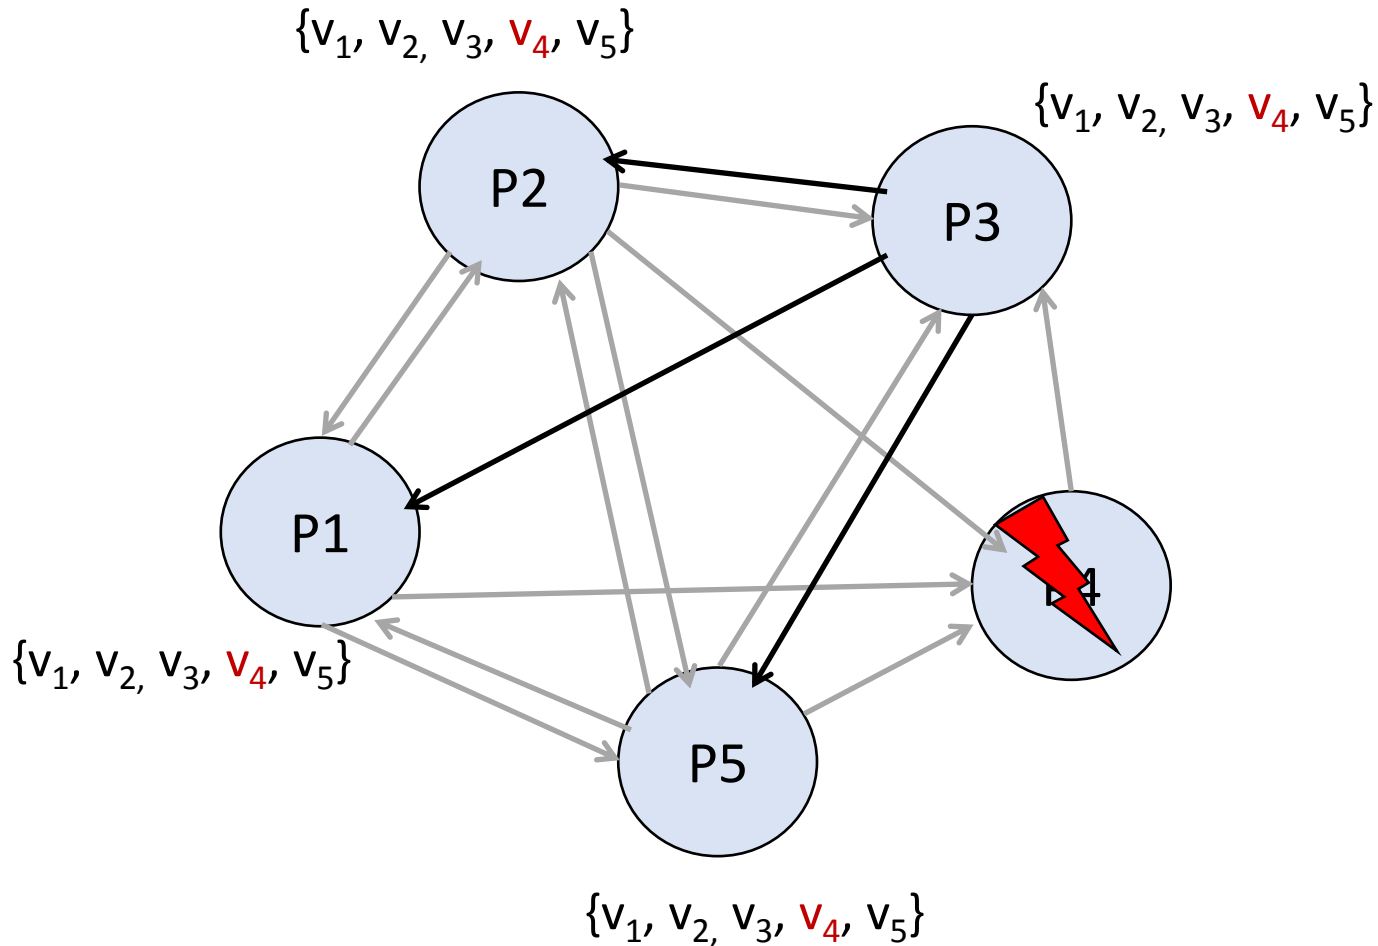
Need R-multicast

B-multicast is not enough for this




Need R-multicast

B-multicast is not enough for this




Need R-multicast

Handling failures




- P4 fails before sending v_4 to anyone.
- What should other processes do?
- Detect failure. *Timeout!*
- Assume proposals are sent at time 's'.
- Worst-case skew is ϵ .
- Maximum message transfer time (including local processing) is T.
- What should timeout value be?

Handling failures




- Assume proposals sent at time s .
- Worst-case skew is ϵ .
- Maximum message transfer time (including local processing) is T .
- What should timeout value be?
- How about $\epsilon + T$?
 - P_i waits for $(\epsilon + T)$ time units after sending proposal at time s
 - Any other process must have sent proposed value before $s + \epsilon$.
 - The proposed value should have reached P_i by $(s + \epsilon + T)$.
 - *Will this work?*

Handling failures




- Assume proposals sent at time s .
- Worst-case skew is ϵ .
- Maximum message transfer time (including local processing) is T .
- What should timeout value be?
- How about $\epsilon + T$?
 - Local time at a process P_i .
 - P_j must have sent proposed value before time $s + \epsilon$.
 - The proposed value should have reached P_i by $(s + \epsilon + T)$.
 - *Will this work?*

Handling failures




- Assume proposals sent at time s .
- Worst-case skew is ϵ .
- Maximum message transfer time (including local processing) is T .
- What should timeout value be?
- How about $\epsilon + 2 \cdot T$?
 - *Will this work?*

Handling failures




- Assume proposals sent at time s .
- Worst-case skew is ϵ .
- Maximum message transfer time (including local processing) is T .
- What should timeout value be?
- How about $\epsilon + 2 \cdot T$?
 - *Will this work?*

Handling failures




- Assume proposals sent at time s .
- Worst-case skew is ϵ .
- Maximum message transfer time (including local processing) is T .
- What should timeout value be?
- How about $\epsilon + 3 \cdot T$?
 - *Will this work?*

Handling failures



- Assume proposals sent at time s .
- Worst-case skew is ϵ .
- Maximum message transfer time (including local processing) is T .
- What should timeout value be?
- How about $\epsilon + 3 \cdot T$?
 - *Will this work?*

Handling failures



- Assume proposals sent at time s .
- Worst-case skew is ϵ .
- Maximum message transfer time (including local processing) is T .
- What should timeout value be?
- Timeout = $\epsilon + (F+1)*T$ for up to F failed process.

Also holds for R-multicast from a single sender.

Round-based algorithm

- For a system with at most F processes crashing
 - All processes are *synchronized* and operate in “rounds” of time.
 - One round of time is equivalent to $\epsilon + T$ units.
 - At each process, the i^{th} round
 - starts at local time $s + (i - 1) * (\epsilon + T)$
 - ends at local time $s + i * (\epsilon + T)$
 - The start or end time of a round in two different processes differs by at most ϵ .
 - The algorithm proceeds in $F + 1$ rounds.
 - Assume communication channels are reliable.

Round-based algorithm

Values^r_i: the set of proposed values known to P_i at the beginning of round r.

Initially Values¹_i = {v_i}

for round = 1 to F+1 do

 B-multicast (Values^r_i – Values^{r-1}_i)

 // iterate through processes, send each a message

 Values^{r+1}_i ← Values^r_i

 wait until one round of time expires.

 for each v_j received in this round

 Values^{r+1}_i = Values^{r+1}_i ∪ v_j

 end

end

d_i = minimum(Values^{F+2}_i)

Why does this work?

- After $F+1$ rounds, all non-faulty processes would have received the same set of values.
- *Proof by contradiction.*
- Assume that two non-faulty processes, say P_i and P_j , differ in their final set of values (i.e., after $F+1$ rounds)
- Assume that P_i possesses a value v that P_j does not possess.
 - P_i must have received v in the **very last** round, else p_i would have sent v to p_j in that last round
 - So, in the last round: a third process, P_k , must have sent v to P_i , but then crashed before sending v to P_j .
 - Similarly, a fourth process sending v in the **last-but-one round** must have crashed; otherwise, both P_k and P_j should have received v .
 - Implies at least one (unique) crash in each of the preceding rounds.
 - This means total of $F+1$ crashes contradicts assumption up to F crashes.

Consensus in synchronous systems

Dolev and Strong proved that for a system with up to F failures (or faulty processes), at least $F+1$ rounds of information exchange is required to reach an agreement.

What about asynchronous systems?

- Using time-based “rounds” or timeouts may not work.
 - Cannot guarantee both completeness and accuracy for failure detection.
 - Cannot differentiate between an extremely slow process and a failed process.
 - Key intuition behind the famous FLP result on the impossibility of consensus in asynchronous systems.
 - *Impossibility of Distributed Consensus with One Faulty Process, Fischer-Lynch-Paterson (FLP), 1985*
 - Stopped distributed system designers dead in their tracks.
 - A lot of claims of “reliability” vanished overnight.
- (Proof is not in your syllabus – optional self-study)*

What about asynchronous systems?

- We cannot “solve” consensus in asynchronous systems.
 - We cannot meet both safety and liveness requirements.
 - Maybe it is ok to guarantee just one requirement.
- Option 1:
 - Set a super conservative timeout for terminating algorithm.
 - Safety violated if a process (or network) is very, very slow.
- Option 2:
 - Let's focus on guaranteeing *safety* under all possible scenarios.
 - If real situation not too dire, hopefully the algorithm terminates.

Paxos Consensus Algorithm

- Paxos algorithm for consensus in asynchronous systems.
 - Most popular consensus-algorithm.
 - A lot of systems use it
 - Zookeeper (Yahoo!), Google Chubby, and many other companies.
 - Not guaranteed to terminate, but never violates safety.

Paxos Consensus Algorithm

- *Guess who invented it?*
 - Leslie Lamport!
- Original paper: The Part-time Parliament.
 - Used analogy of a “part-time parliament” on an ancient Greek island of Paxos.
 - No one understood it.
 - The paper was rejected.
- Published “*Paxos made simple*” 10 years later.


Paxos Algorithm

- Three types of roles:
 - Proposers: propose values to *acceptors*.
 - All or subset of processes.
 - Having a *single proposer* (leader) may allow faster termination.
 - Acceptors: accept proposed values (under certain conditions).
 - All or subset of processes.
 - Learners: learns the value that has been accepted by *majority* of acceptors.
 - All processes.

Paxos Algorithm: Try 1: Single Phase

- A proposer multicasts its proposed value to a large enough set (larger than majority) of acceptors.
- An acceptor accepts the first proposed value it receives.
- If majority of acceptors have accepted the same value v , then v is the decided value.
- *What can go wrong here?*

Paxos Algorithm: Try 1: Single Phase




No decision reached!

Paxos Algorithm: Proposal numbers

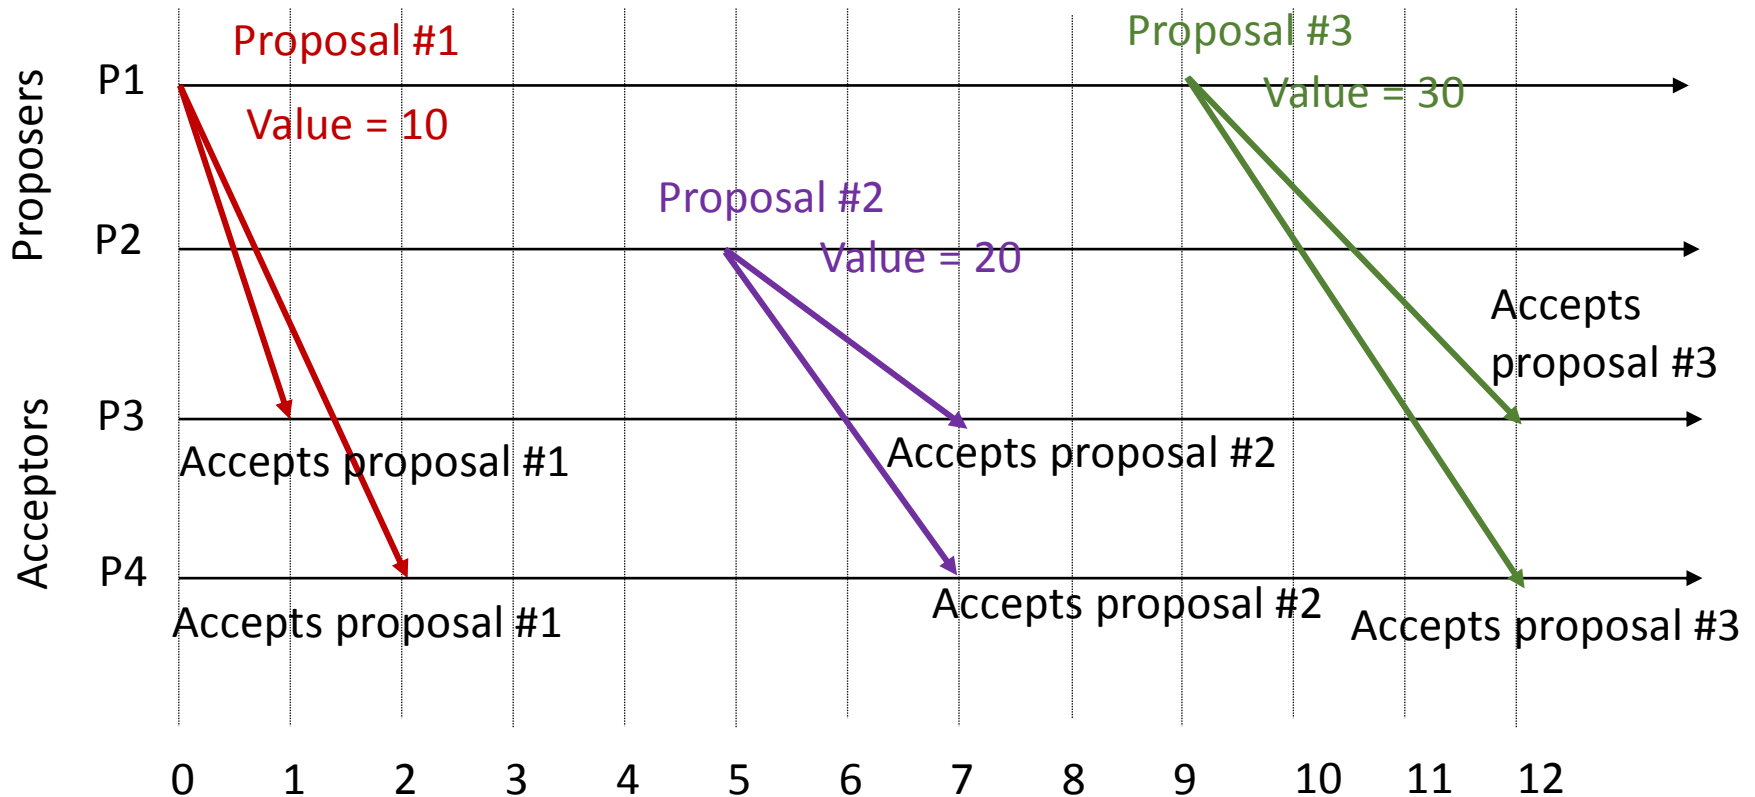
- Allow an acceptor to accept multiple proposals.
 - Accepting is different from *deciding*.
- Distinguish proposals by assigning unique ids (a **proposal number**) to each proposal.
 - Configure a disjoint set of possible proposal numbers for different processes.
 - Proposal number is different from proposed value!
- A higher number proposal overwrites and pre-empts a lower number proposal.

Paxos Algorithm: Try 2: Proposal #s



What can go wrong here?

Paxos Algorithm: Try 2: Proposal #s

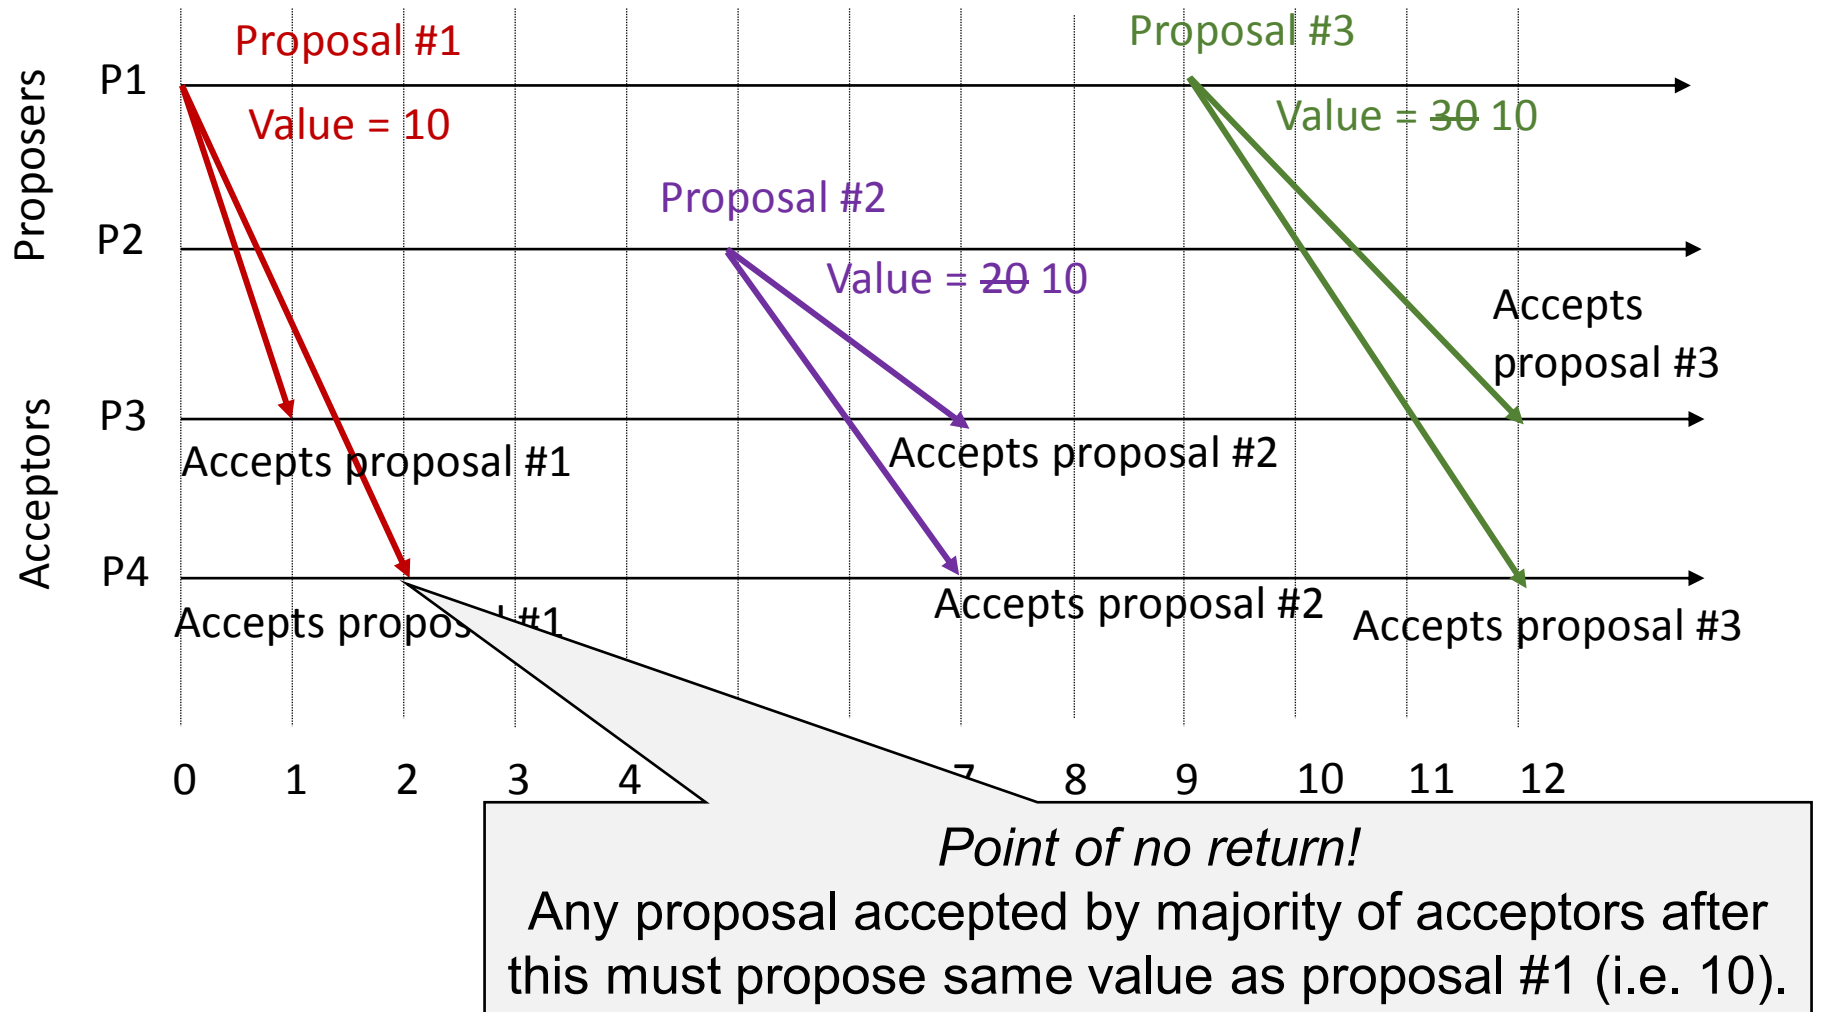


When do we stop and decide on a value?

Paxos Algorithm

- Key condition:
 - When majority of acceptors accept a single proposal with a value v , then that value v becomes the decided value.
 - This is an implicit decision. Learners may not know about it right-away.
 - Any higher-numbered proposal that gets accepted by majority of acceptors after the implicit decision must propose the same decided value.

Paxos Algorithm



Paxos Algorithm: Two phases

- Phase 1:
 - A proposer selects proposal number (n), sends **prepare** request with n to majority of acceptors requesting:
 - Promise me you will not reply to any other proposal with a lower number.
 - Promise me you will not accept any other proposal with a lower number.
 - If an acceptor receives a **prepare** request for proposal $\#n$, and it has not responded to a **prepare** request with a higher number, it replies back saying:
 - **OK!** I will make that promise for any request I receive in future.
 - (If applicable) I have already accepted a value v from a proposal with lower number $m < n$. The proposal has the highest number among the ones I accepted so far.

Paxos Algorithm: Two phases

- Phase 2:
 - If a proposer receives an **OK** response for its **prepare** request $\#n$ from a *majority* of acceptors, then it sends an **accept** request with a proposed value. What is the proposed value?
 - The value v of the *highest numbered proposal* among the received responses.
 - Any value if no previously accepted value in the received responses.
 - If an acceptor receives an **accept** request for proposal $\#n$, and it has not responded a **prepare** request with a higher number, it **accepts** the proposal.

Next Class

- Wrap up discussion on Paxos algorithm
 - Why it guarantees safety?
 - How do processes learn about the decided value.
- Raft: Log-based consensus

Summary

- Consensus is a fundamental problem in distributed systems.
- Consensus possible in synchronous systems.
 - Algorithm based on time-synchronized rounds.
 - Need at least $(F+1)$ rounds to handle up to F failures.
- Consensus impossible in asynchronous systems.
 - Cannot distinguish between timeout and a very slow process.
 - Paxos algorithm:
 - Guarantees safety but not liveness.
 - Hopes to terminate if under good enough conditions.