

Distributed Systems

ECE428

Lecture 16

Adopted from Spring 2021

Recap

- Log consensus / Replicated State Machines
 - Need to ensure the log entries / commands are processed in the same order at different servers.
 - Same requirement as total reliable multicast.
- Paxos is the most popular consensus algorithm for asynchronous system. Multi-paxos for log consensus.
 - Difficult to understand and implement.
- Raft designed for simplicity.

Recap: Raft

- Elect a leader. Only leader can add and commit log entries.
 - Tries to eventually make other server's logs identical to its own.
- Handling failures: each new leader election increments the *term*. Terms increase monotonically over time.
- Raft design provides various forms of guarantees:
 - Only one leader elected per term
 - A server can only grant one vote per term. Need majority to win.
 - If log entries on different servers have same index and term, they store the same command, and the logs are identical in all preceding entries.
 - Consistency check when appending a new entry.
 - If a leader has decided that a log entry is committed, that entry will be present in the logs of all future leaders.
 - Restrictions on commitment and leader election.
- Together ensure: *Once log entry has been applied to state machine, no other state machine can apply a different value for that log entry.*

Agenda for today

- Consensus

- Consensus in synchronous systems
 - *Chapter 15.4*
- Impossibility of consensus in asynchronous systems
 - *We will not cover the proof in details*
- Good enough consensus algorithm for asynchronous systems:
 - *Paxos made simple, Leslie Lamport, 2001*
- Other forms of consensus algorithm
 - Raft (log-based consensus)
 - Block-chains / Bitcoins (distributed consensus)

Bitcoins

- Implement a *distributed* replicated state machine that maintains an *account ledger* (= *bank*).
 - No user should be able to “double-spend”.
 - Need to know about all transactions to validate this.
 - Who does this validation? Cannot trust a single central authority.
 - Any participant (replica) should be able to validate.
 - All replicas must agree on the single history of transactions and their ordering.
- Scale to thousands of replicas distributed across the world.
- Allow old replicas to fail, new replicas to join seamlessly.
- Must withstand various types of attacks.

Uses Blockchains for Consensus

- Why not use Paxos / Raft?
 - Need to scale to thousands of replicas across the world.
 - May not even know of all replicas a priori.
 - Participants may leave / join dynamically.
 - Paxos/Raft are ill-suited for such a setup.
 - Leader election in Raft or proposals in Paxos require communication with at least a majority of servers.
 - Require knowing the number of replicas.
 -
- *So how does blockchain work?*
 - Focus of today's class. Only a high-level discussion.

Basic Idea

Transactions grouped into a *block* that gets added to the *chain* (history of transactions) by the “leader of that block”.

Lottery Leader Election

- Every node chooses a random number
- Leader = “closest to 0”

Lottery Leader Election

- Every node chooses a random number
 - The method for choosing the number in blockchains enables log consensus (with a high probability).
 - Requires the leader to expend CPU (as *proof-of-work*).
- Leader = “closest to 0”
 - Defined such that a replica can determine this independently without coordination

Choosing the random number

- Cryptographic hash function:
 - $H(x) \rightarrow \{0, 1, \dots, 2^{256}-1\}$
- Hard to *invert*:
 - Given y , find x such that $H(x) = y$
 - E.g., SHA256, SHA3, ...
- Every node picks random number x and computes $H(x)$
- Node with $H(x)$ “closest to 0” wins
 - Finding such an x requires expending CPU (*proof-of-work*).
- *But once we have found an ‘x’, we can always be the leader for all blocks, or even share it with colluding parties. How to prevent that?*

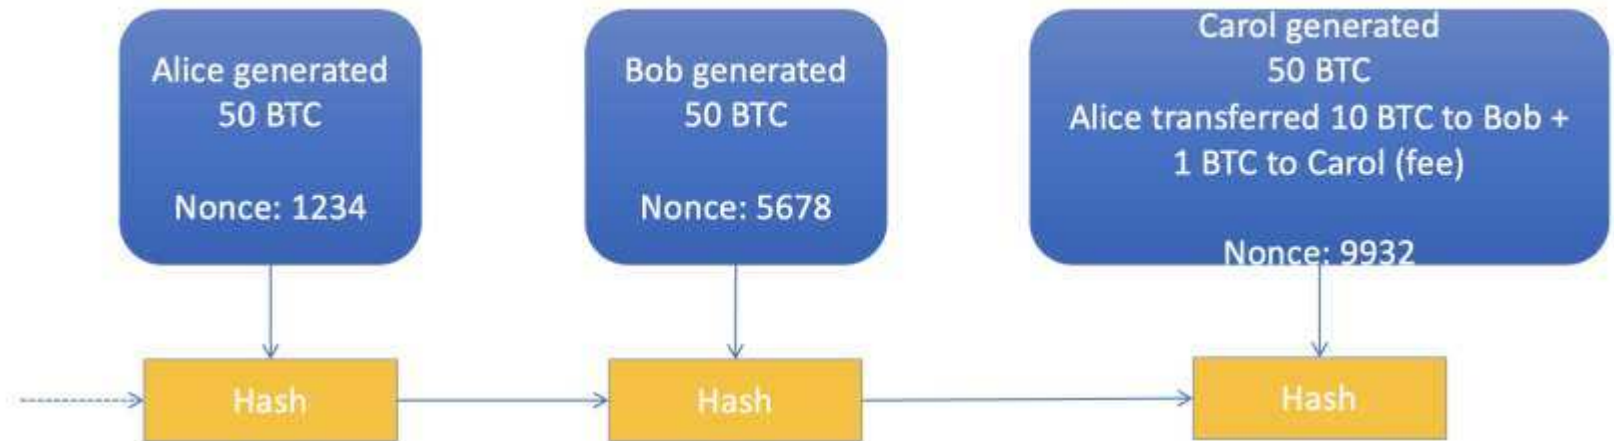
Using a seed

- Every node picks x , computes $H(\text{seed} || x)$
- Closest to 0 wins
- What to use as a seed?
- Hash of:
 - Previous log
 - Node identifier
 - New messages to add to log
- *How to find “closest to 0”?*

Iterated Hashing / Proof of work

- Repeat:
 - Pick random x , compute $y = H(\text{seed} \parallel x)$
 - If $y < T$, you win!
- Set threshold T so that on average, one winner every few minutes
- Given a *solution*, x such that $H(\text{seed} \parallel x) < T$, anyone can *verify* the solution in constant time (microseconds).

Chaining the blocks



Account	Balance
Alice	39 BTC
Bob	60 BTC
Carol	51 BTC

Protocol Overview


- New transactions broadcast to all nodes.
- Each node collects new transactions into a block.
- Each node works on finding a proof-of-work for its block to become its leader and get it appended to a chain.
 - i.e. finds x , such that $H(\text{seed} || x) < T$.
- When a node finds a proof-of-work, it broadcasts it to all nodes.
- Nodes accept a block only if all transactions in it are valid.
- Nodes express their acceptance by working on creating the next block in the chain, using the hash of accepted block as previous hash.

What could go wrong?

- Two nodes may end up mining different versions of next block.

Longest Chain Rule

- Two nodes may end up mining different versions of next block.
- A node may receive two versions of the next block.
- Will store both, but work on the first one they receive.
- Over time, more blocks will be received.
- The node will switch to working on the *longest chain*.



Security Property

- Majority decision is represented by the longest chain.
 - It has greatest “proof-of-work” invested in it.
- If majority of CPU power is controlled by honest nodes, the honest chain will grow fastest and outpace competing chains.
- To modify past blocks, an attacker will need to redo the proof-of-work for that block, and all blocks after it, and then surpass the work of honest nodes.
- Probability of attack reduces as more blocks get added.

Incentives for Logging

- Security better if more people participated in logging.
- Incentivize users to log *others'* transactions
 - Transaction fees: pay me x% to log your data
 - Mining reward: each block *creates* bitcoins

Logging Speed

- How to set T ?
 - Too large waste effort due to broadcast delays & chain splits
 - Too small: slows down transactions
- Periodically adjust difficulty T such that one block gets added every 10 minutes.
 - Depends on hardware speed (which increases over the years) and number of participants (which vary over time).
- Determined algorithmically based on measured average number of blocks mined per hour.

Bitcoin Broadcast

- Need to broadcast:
 - Transactions to all nodes, so they can be included in a block.
 - New blocks to all nodes, so that they can switch to longest chain.
- What if we use R-multicast?
 - Have to send $O(N)$ messages
 - Have to *know* which nodes to send to
 - Not a suitable choice.

Gossip / Viral propagation

- Each node connects to a small set of *neighbors* (10–100).
- Nodes propagate transactions and blocks to neighbors.
- Push method: when you hear a new transaction/block, resend them to all (some) of your neighbors (flooding).
- Pull method: periodically poll neighbors for list of blocks/transactions, then request any you are missing.
- Unreliable: some nodes may not receive all transactions or all blocks. But that's ok.

Maintaining Neighbors

- A *seed* service
 - Gives out a list of random or well-connected nodes
 - E.g., seed.bitnodes.io
- Neighbor discovery
 - Ask neighbors about *their* neighbors
 - Randomly connect to some of them

Bitcoin Summary

- Unreliable broadcast using gossip
- Probabilistic “leader” election for mining blocks (tx ordering)
- Longest chain rule to ensure long-term (probabilistic) consistency and security
- Compared with Paxos/Raft:
 - Scales to thousands of participants, dynamic groups
 - Tens of minutes to successfully log a transaction (vs. milliseconds)